

x0 Protocol: A Decentralized Payment Infrastructure for Autonomous Agents

Technical Whitepaper v1.0

Blessed Tosin-Oyinbo Olamide
x0protocol.dev

February 2026

Abstract

We present x0, a decentralized protocol for autonomous agent payments built on Solana. The protocol introduces a spending policy enforcement mechanism using Token-2022 transfer hooks, enabling programmable spending limits, whitelist verification, and privacy controls for AI agents. x0 implements HTTP 402 (Payment Required) for standardized payment negotiation, conditional escrow with dispute resolution, on-chain reputation scoring with temporal decay, a USDC-backed wrapper token with cryptographic reserve invariants, and on-chain zero-knowledge proof verification for confidential transfers using Groth16 proofs over the Ristretto255 curve. We further introduce FROSTGATE, a trustless bidirectional Base \longleftrightarrow Solana bridge that combines Hyperlane message passing with SP1 STARK proofs to enable cross-chain agent payments without trusted intermediaries. The system achieves trustless agent-to-agent transactions while preserving human oversight through a “Blink” mechanism for exceptional cases. We formally verify the protocol’s security properties and demonstrate its efficiency through cryptographic proofs and empirical benchmarks.

Contents

1	Introduction	6
1.1	Motivation	6
1.2	Contributions	6
1.2.1	Architecture Overview	8
2	Preliminaries	8
2.1	Cryptographic Primitives	8
2.2	Solana Architecture	9
2.2.1	Account Model	9
2.2.2	Program Derived Addresses (PDAs)	9
2.2.3	Token-2022 Extensions	9
2.3	Bloom Filters	9
2.4	Merkle Trees	10

3 Policy Enforcement Layer (x0-guard)	10
3.1 Problem Statement	10
3.2 Design Goals	10
3.3 Architecture	10
3.3.1 AgentPolicy Account	10
3.3.2 Transfer Hook Mechanism	11
3.4 Rolling Window Algorithm	11
3.5 Whitelist Verification	12
3.5.1 Merkle Mode	12
3.5.2 Bloom Mode	12
3.5.3 Domain Mode	12
3.6 Privacy Levels	13
3.6.1 Public Mode	13
3.6.2 Confidential Mode	13
3.7 Reentrancy Protection	13
4 HTTP 402 Protocol	14
4.1 Problem Statement	14
4.2 Protocol Flow	14
4.3 Protocol Design	14
4.3.1 Payment Request	14
4.3.2 Payment Proof	15
4.3.3 Verification	15
4.4 Security Properties	16
5 Conditional Escrow (x0-escrow)	16
5.1 Design	16
5.1.1 Escrow State Machine	16
5.1.2 Escrow Account	16
5.2 Key Operations	17
5.2.1 Create and Fund	17
5.2.2 Delivery and Release	17
5.2.3 Auto-Release	17
5.2.4 Dispute Resolution	19
5.3 Security Analysis	19
6 Reputation Oracle (x0-reputation)	19
6.1 Design Goals	19
6.2 Reputation Account	19
6.3 Reputation Score Calculation	20
6.4 Temporal Decay	21
6.5 Authorization Model	21

7 Agent Discovery (x0-registry)	21
7.1 Registry Entry	21
7.2 Discovery Algorithm	22
7.3 Capability Metadata	22
8 USDC Wrapper (x0-wrapper)	23
8.1 Problem Statement	23
8.2 Design	23
8.2.1 State Accounts	23
8.2.2 Reserve Invariant	24
8.3 Governance	24
8.3.1 Timelock	24
8.3.2 Emergency Pause	25
8.4 Reserve Ratio Monitoring	25
9 Human-in-the-Loop (Blinks)	25
9.1 Problem Statement	25
9.2 Blink Design	25
9.2.1 Generation	26
9.2.2 Rate Limiting	26
9.2.3 Approval	26
10 Economic Model	26
10.1 Fee Structure	26
10.2 Fee Collection	26
10.3 Token Economics	27
10.3.1 x0-USD Supply Dynamics	27
11 Zero-Knowledge Proof Verification (x0-zk-verifier)	27
11.1 Problem Statement	27
11.2 Cryptographic Foundations	28
11.2.1 Twisted ElGamal Encryption	28
11.2.2 Groth16 Proof System	28
11.3 Proof Types	28
11.3.1 PubkeyValidityProof	29
11.3.2 ZeroBalanceProof	29
11.3.3 WithdrawProof	29
11.3.4 TransferProof	29
11.4 Architecture	29
11.4.1 Proof Context Account	29
11.4.2 Proof Freshness	30
11.4.3 Amount Bounds	30
11.5 Off-Chain Proof Generation (x0-zk-proofs)	30
11.6 Verification Flow	31
11.7 Security Properties	31

12 Security Analysis	31
12.1 Threat Model	31
12.2 Attack Scenarios and Mitigations	32
12.2.1 Compromised Agent Key	32
12.2.2 Sybil Attack on Reputation	32
12.2.3 Escrow Griefing	32
12.2.4 Wrapper Reserve Drain	32
12.2.5 Transfer Hook Configuration Attack	33
12.3 Formal Verification	33
12.3.1 Spend Limit Invariant	33
12.3.2 Reserve Invariant	33
13 Performance Analysis	33
13.1 Computational Complexity	33
13.2 On-Chain Costs	33
13.3 Transaction Benchmarks	33
14 Related Work	33
14.1 Payment Channels	33
14.2 Escrow Protocols	35
14.3 Reputation Systems	35
15 FROSTGATE: Trustless Cross-Chain Bridge	35
15.1 Problem Statement	35
15.2 Definitions	36
15.3 Architecture Overview	36
15.4 On-Chain State	37
15.4.1 Bridge Configuration	37
15.4.2 Bridge Message	37
15.4.3 EVM Proof Context	37
15.4.4 Bridge Out Message	38
15.5 Inbound Protocol (Base → Solana)	38
15.5.1 Step 1: USDC Lock on Base	38
15.5.2 Step 2: Hyperlane Message Delivery	39
15.5.3 Step 3: STARK Proof Verification	39
15.5.4 Step 4: Mint Execution	40
15.6 Outbound Protocol (Solana → Base)	40
15.6.1 Step 1: Burn on Solana	40
15.6.2 Step 2: Proof Generation and Unlock on Base	42
15.7 SP1 STARK Proof Circuits	42
15.7.1 EVM Proof Circuit (Base → Solana)	42
15.7.2 Solana Proof Circuit (Solana → Base)	43
15.8 Rate Limiting and Circuit Breakers	44
15.8.1 Per-Transaction Bounds	44
15.8.2 Daily Rolling Window	44

15.8.3 Circuit Breakers	44
15.9 Proof Freshness	45
15.10EVM Contracts	45
15.10.1 X0LockContract	45
15.10.2 X0UnlockContract	45
15.11Administrative Controls	46
15.12Security Analysis	46
15.12.1 Threat Model	46
15.12.2 Mitigations	46
16 Future Work	48
16.1 Extended Zero-Knowledge Applications	48
16.2 Cross-Chain Extension	48
16.3 Machine Learning Integration	49
16.4 Governance	49
17 Conclusion	49
A Error Codes Reference	51
B Program Addresses	53
C Mathematical Notation	53
D SDK Integration Guide	53
D.1 Installation	53
D.2 Quick Start	54
D.3 SDK Modules	54
D.4 Resources	54
D.5 Confidential Transfer Architecture	55

1 Introduction

1.1 Motivation

The proliferation of autonomous AI agents in production environments necessitates robust payment infrastructure. Traditional payment systems are not well suited for agent-to-agent transactions due to:

1. **High latency:** Credit card settlements take 2-3 days, incompatible with real-time agent interactions
2. **High fees:** 2-3% credit card fees are prohibitive for micropayments
3. **Gatekeeping:** Centralized payment processors can deny service arbitrarily
4. **Lack of programmability:** Traditional payments cannot enforce complex spending rules
5. **Privacy concerns:** All transaction data visible to payment processors

Blockchain-based payments solve latency and gatekeeping but introduce new challenges:

1. **Custody risk:** Agents require private keys for signing, creating attack surface
2. **Unbounded spending:** Standard wallets lack programmable spending limits
3. **No recourse:** Blockchain transactions are irreversible by design
4. **Discovery problem:** Finding trustworthy service providers in decentralized systems

1.2 Contributions

We present x_0 , which makes the following contributions:

1. **Programmable Spending Policies:** A cryptographic policy enforcement layer using Solana's Token-2022 transfer hooks that validates every transaction against owner-defined rules
2. **HTTP 402 Protocol:** An extension to HTTP status codes enabling standardized payment negotiation between agents and services
3. **Conditional Escrow:** A trustless escrow mechanism with optional third-party arbitration for high-value transactions
4. **Reputation Oracle:** An on-chain reputation system with temporal decay, preventing stale reputations from dominating
5. **USDC Wrapper:** A 1:1 USDC-backed token with cryptographic reserve invariants and timelocked governance
6. **Human-in-the-Loop (Blinks):** A fallback mechanism for exceptional transactions requiring human approval

7. **Zero-Knowledge Verification:** On-chain Groth16 proof verification for confidential transfers, with client-side proof generation compiled to WebAssembly
8. **FROSTGATE Cross-Chain Bridge:** A trustless bidirectional Base \longleftrightarrow Solana bridge using Hyperlane message passing and SP1 STARK proofs, with event-level validation, rate limiting, and circuit breakers

Contributions 1–7 form the core x0 protocol, operating entirely on Solana. Contribution 8 (FROSTGATE) is an optional extension enabling trustless cross-chain interactions.

The x0 protocol consists of nine interoperating programs deployed on Solana, two EVM contracts on Base, two SP1 STARK proof circuits, and an off-chain WASM cryptographic library:

- **x0-guard:** Policy enforcement via transfer hooks
- **x0-token:** Token-2022 mint configuration with confidential transfer support
- **x0-escrow:** Conditional payment escrow with dispute resolution
- **x0-registry:** Agent discovery and capability advertisement
- **x0-reputation:** Trust scoring and transaction history with temporal decay
- **x0-wrapper:** USDC-backed stable wrapper token with timelocked governance
- **x0-zk-verifier:** On-chain Groth16 zero-knowledge proof verification for confidential transfers
- **x0-bridge:** Cross-chain bridge program (FROSTGATE) with Hyperlane integration and SP1 STARK proof verification
- **x0-common:** Shared types, constants, error codes, and utilities

The protocol also includes EVM and off-chain components:

- **X0LockContract** (Solidity): USDC lock contract on Base for inbound bridge transfers
- **X0UnlockContract** (Solidity): USDC release contract on Base for outbound bridge transfers
- **sp1-evm-prover:** SP1 STARK circuit proving EVM transaction inclusion via MPT proofs
- **sp1-solana-prover:** SP1 STARK circuit proving Solana account inclusion via fanout-16 Merkle proofs and validator quorum
- **x0-zk-proofs:** WebAssembly module compiled from Rust, providing client-side Groth16 proof generation using `solana-zk-token-sdk`

1.2.1 Architecture Overview

The x0 protocol is organized in four layers. At the top, external Solana ecosystem programs—USDC (SPL Token) and Token-2022—provide the underlying token infrastructure and extension framework. The core layer comprises nine on-chain programs: **x0-token** configures the Token-2022 mint with transfer hook and confidential transfer extensions; **x0-guard** implements the transfer hook, enforcing per-transaction and daily spending policies via Merkle or Bloom filter whitelists and invoking **x0-reputation** via CPI for trust-aware validation; **x0-wrapper** manages USDC-to-x0-USD wrapping with cryptographic reserve invariants; **x0-escrow** provides conditional payment escrow with dispute resolution, also backed by reputation CPI; **x0-registry** handles agent discovery and capability advertisement; **x0-zk-verifier** performs on-chain Groth16 proof verification for confidential transfers; and **x0-bridge** orchestrates cross-chain bridging via Hyperlane message passing and SP1 STARK proof verification, with CPI into x0-wrapper for reserve-backed minting and burning. All on-chain programs share types, constants, and error codes through the **x0-common** library crate. In the off-chain layer, **x0-zk-proofs** is a Rust-to-WASM module that generates client-side Groth16 proofs for the ZK verifier, while **sp1-evm-prover** and **sp1-solana-prover** are SP1 guest circuits that produce STARK proofs of EVM transaction inclusion and Solana account inclusion, respectively. The cross-chain layer comprises two Solidity contracts on Base—**X0LockContract** (locking USDC and dispatching Hyperlane messages to x0-bridge) and **X0UnlockContract** (verifying SP1 proofs and releasing USDC)—completing the bidirectional bridge.

2 Preliminaries

2.1 Cryptographic Primitives

Definition 2.1 (Hash Function). A cryptographic hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^{256}$ satisfies:

1. **Preimage resistance:** Given h , it is computationally infeasible to find m such that $H(m) = h$
2. **Second preimage resistance:** Given m_1 , it is computationally infeasible to find $m_2 \neq m_1$ such that $H(m_1) = H(m_2)$
3. **Collision resistance:** It is computationally infeasible to find $m_1 \neq m_2$ such that $H(m_1) = H(m_2)$

We use SHA-256 for all hash operations, providing 128-bit security against collision attacks.

Definition 2.2 (Digital Signature Scheme). A signature scheme (**KeyGen**, **Sign**, **Verify**) consists of:

- **KeyGen**(1^λ) $\rightarrow (sk, pk)$: Generates a key pair
- **Sign**(sk, m) $\rightarrow \sigma$: Signs message m with secret key sk
- **Verify**(pk, m, σ) $\rightarrow \{0, 1\}$: Verifies signature σ on message m

Solana uses Ed25519 signatures, providing 128-bit security with 64-byte signatures.

2.2 Solana Architecture

2.2.1 Account Model

Solana uses an account-based model where each account has:

- **Address:** A 32-byte Ed25519 public key
- **Lamports:** Balance in lamports ($1 \text{ SOL} = 10^9 \text{ lamports}$)
- **Data:** Arbitrary byte array storing account state
- **Owner:** The program with write access to the account
- **Executable:** Whether the account contains program code

2.2.2 Program Derived Addresses (PDAs)

A PDA is a deterministic address derived from:

$$\text{PDA}(\text{seeds}, \text{program_id}) = \text{FindProgramAddress}(\text{seeds}, \text{program_id}) \quad (1)$$

Where $\text{FindProgramAddress}$ searches for a public key P such that:

$$P = H(\text{seeds} \parallel \text{program_id} \parallel [b]) \quad \text{and} \quad P \notin E(\mathbb{F}_p) \quad (2)$$

Here $E(\mathbb{F}_p)$ is the Ed25519 elliptic curve, and $b \in [0, 255]$ is the "bump" seed. PDAs are off-curve points, ensuring no private key exists.

2.2.3 Token-2022 Extensions

Token-2022 is Solana's next-generation token program supporting extensions:

- **Transfer Hook:** Calls a specified program on every transfer
- **Transfer Fee:** Withholds a percentage of each transfer
- **Confidential Transfer:** Encrypts balances using ElGamal encryption

2.3 Bloom Filters

Definition 2.3 (Bloom Filter). A Bloom filter is a probabilistic data structure for set membership testing. For a set $S = \{s_1, \dots, s_n\}$:

- Bit array: $B[0..m - 1]$ initialized to 0
- Hash functions: $h_1, \dots, h_k : \{0, 1\}^* \rightarrow [0, m - 1]$
- Insert: For each $s \in S$, set $B[h_i(s)] = 1$ for $i = 1, \dots, k$
- Query: Element $x \in S$ if $B[h_i(x)] = 1$ for all i

Theorem 2.1 (Bloom Filter False Positive Rate). *For a Bloom filter with m bits, n elements, and k hash functions:*

$$P(\text{false positive}) = \left(1 - e^{-kn/m}\right)^k \quad (3)$$

Optimal k minimizes false positives:

$$k_{\text{opt}} = \frac{m}{n} \ln 2 \quad (4)$$

2.4 Merkle Trees

Definition 2.4 (Merkle Tree). A Merkle tree is a binary tree where:

- Leaves: $L_i = H(\text{data}_i)$ for $i = 0, \dots, n - 1$
- Internal nodes: $N_{i,j} = H(N_{i,2j} \parallel N_{i,2j+1})$
- Root: $r = N_{0,0}$

Theorem 2.2 (Merkle Proof Size). *For a tree with n leaves, a membership proof requires $O(\log n)$ hashes.*

3 Policy Enforcement Layer (x0-guard)

3.1 Problem Statement

Consider an agent A owned by user U with private key sk_U . The agent requires a signing key sk_A to perform transactions autonomously. Without constraints, compromise of sk_A allows unbounded spending.

3.2 Design Goals

1. **Spend Limits:** Enforce maximum spending in rolling 24-hour windows
2. **Transaction Limits:** Cap individual transaction sizes
3. **Whitelist Verification:** Restrict recipients to approved addresses
4. **Privacy:** Support confidential (encrypted) transfers
5. **Auditability:** Maintain on-chain transaction history
6. **Revocability:** Allow owner to revoke agent authority

3.3 Architecture

3.3.1 AgentPolicy Account

Each agent has a Program Derived Address (PDA) storing its policy:

```

1 #[account]
2 pub struct AgentPolicy {
3     pub version: u8,           // Account version (migration)
4     pub owner: Pubkey,        // Cold wallet (full control)
5     pub agent_signer: Pubkey, // Hot key (delegated)
6     pub daily_limit: u64,      // Max spend per 24h
7     pub max_single_transaction: Option<u64>,
8     pub rolling_window: Vec<SpendingEntry>,
9     pub privacy_level: PrivacyLevel,
```

```

10  pub whitelist_mode: WhitelistMode,
11  pub whitelist_data: WhitelistData,
12  pub is_active: bool,
13  pub require_delegation: bool,    // Require token delegation
14  pub bound_token_account: Option<Pubkey>,
15  pub last_update_slot: u64,      // For rate limiting
16  pub auditor_key: Option<Pubkey>, // Optional auditor
17  pub blink_hour_start: i64,      // Blink rate limit window
18  pub blinks_this_hour: u8,       // Blinks in current window
19  pub bump: u8,
20 }
21
22 pub struct SpendingEntry {
23     pub amount: u64,
24     pub timestamp: i64,
25 }
```

Listing 1: AgentPolicy Account Structure

3.3.2 Transfer Hook Mechanism

Token-2022's transfer hook enables us to intercept every transfer. The flow is:

1. User initiates transfer of x tokens to recipient R
2. Token-2022 calls x0-guard's `validate_transfer`
3. x0-guard verifies:
 - Signer is authorized agent: `signer = policy.agent_signer`
 - Spend limit not exceeded: $\sum_{t > t_{\text{now}} - 86400} \text{amount}_t + x \leq \text{daily_limit}$
 - Transaction limit not exceeded: $x \leq \text{max_single_transaction}$
 - Recipient whitelisted: $R \in W$ (if whitelist enabled)
4. If validation passes, transfer proceeds; otherwise, reverts

3.4 Rolling Window Algorithm

Theorem 3.1 (Rolling Window Correctness). *For a daily limit L and current time t , the rolling window algorithm ensures:*

$$\sum_{i: t_i > t - 86400} x_i \leq L \quad (5)$$

at all times t .

Proof. By induction on transactions. Base case: initially, the sum is $0 \leq L$. Inductive step: assume the invariant holds before transaction j with amount x_j . The algorithm rejects if:

$$\sum_{i: t_i > t_j - 86400} x_i + x_j > L \quad (6)$$

Therefore, if accepted:

$$\sum_{i: t_i > t_j - 86400} x_i + x_j \leq L \quad (7)$$

Algorithm 1 Rolling Window Spend Limit Enforcement

```
1: procedure VALIDATETRANSFER(policy,  $x, t_{\text{now}}$ )
2:    $t_{\text{cutoff}} \leftarrow t_{\text{now}} - 86400$                                  $\triangleright 24 \text{ hours ago}$ 
3:   policy.rolling_window  $\leftarrow$  policy.rolling_window.retain( $|e| e.ttimestamp > t_{\text{cutoff}}$ )
4:   current_spend  $\leftarrow \sum_{e \in \text{rolling\_window}} e.\text{amount}$ 
5:   if current_spend +  $x > \text{policy}.\text{daily\_limit}$  then
6:     return Error::DailyLimitExceeded
7:   end if
8:   if  $|\text{policy}.\text{rolling\_window}| \geq \text{MAX\_ENTRIES}$  then
9:     return Error::WindowOverflow
10:  end if
11:  policy.rolling_window.push({amount :  $x, \text{timestamp} : t_{\text{now}}$ })
12:  return Success
13: end procedure
```

After adding x_j to the window, the sum equals $\sum_{i: t_i > t_j - 86400} x_i + x_j \leq L$. \square \square

3.5 Whitelist Verification

Three whitelist modes are supported:

3.5.1 Merkle Mode

Store Merkle root r in policy. For transfer to R :

1. Agent provides proof $\pi = \{h_1, \dots, h_{\log n}\}$
2. Verify: $\text{ComputeRoot}(H(R), \pi) = r$

Advantages: $O(\log n)$ proof size, deterministic verification

Disadvantages: Proof must be provided by agent, updates require new root

3.5.2 Bloom Mode

Store Bloom filter B in policy. For transfer to R :

1. Compute $h_i(R)$ for $i = 1, \dots, k$
2. Check: $B[h_i(R)] = 1$ for all i

Advantages: $O(1)$ verification, no proof required

Disadvantages: False positives, filter stored on-chain (4KB)

For $n = 1000$ addresses, $m = 4096 \times 8 = 32768$ bits, $k = 7$:

$$P(\text{FP}) = \left(1 - e^{-7 \times 1000 / 32768}\right)^7 \approx 0.008 = 0.8\% \quad (8)$$

3.5.3 Domain Mode

Store domain prefixes $\{d_1, \dots, d_m\}$ (first 8 bytes of addresses). For transfer to R :

1. Extract prefix: $p = R[0..7]$

2. Check: $p \in \{d_1, \dots, d_m\}$

Advantages: Allows "vanity addresses", compact storage

Disadvantages: Lower security (8-byte prefixes), linear scan

3.6 Privacy Levels

3.6.1 Public Mode

Standard SPL transfers with visible amounts.

3.6.2 Confidential Mode

Uses Token-2022 confidential transfer extension:

1. Balances encrypted with ElGamal: $C = (g^r, g^r \cdot h^b)$ where b is balance
2. Transfers use range proofs to prevent overflow
3. Optional auditor can decrypt amounts

Definition 3.1 (ElGamal Encryption). Public key: (g, h) where $h = g^x$ and x is secret. To encrypt m :

$$\text{Enc}(m) = (g^r, h^r \cdot g^m) \quad (9)$$

for random r . Decryption:

$$\text{Dec}((c_1, c_2)) = \frac{c_2}{c_1^x} = g^m \quad (10)$$

Then solve discrete log to recover m (feasible for small m).

3.7 Reentrancy Protection

Theorem 3.2 (State-Before-Transfer Invariant). *For all escrow/wrapper operations, state updates occur before token transfers. This prevents reentrancy attacks.*

```

1 pub fn release_funds(ctx: Context<ReleaseFunds>) -> Result<()> {
2     let escrow = &mut ctx.accounts.escrow;
3
4     // CRITICAL: Update state BEFORE transfer
5     let amount = escrow.amount;
6     escrow.state = EscrowState::Released;
7
8     // Now transfer (if reentrant call occurs, state check fails)
9     token::transfer(/* ... */, amount)?;
10
11     Ok(())
12 }
```

Listing 2: Reentrancy Protection Pattern

4 HTTP 402 Protocol

4.1 Problem Statement

Existing payment protocols lack standardization for agent-to-agent negotiation. HTTP provides status codes for various conditions but lacks a payment-specific code beyond 402 (Payment Required), which was reserved but never specified.

4.2 Protocol Flow

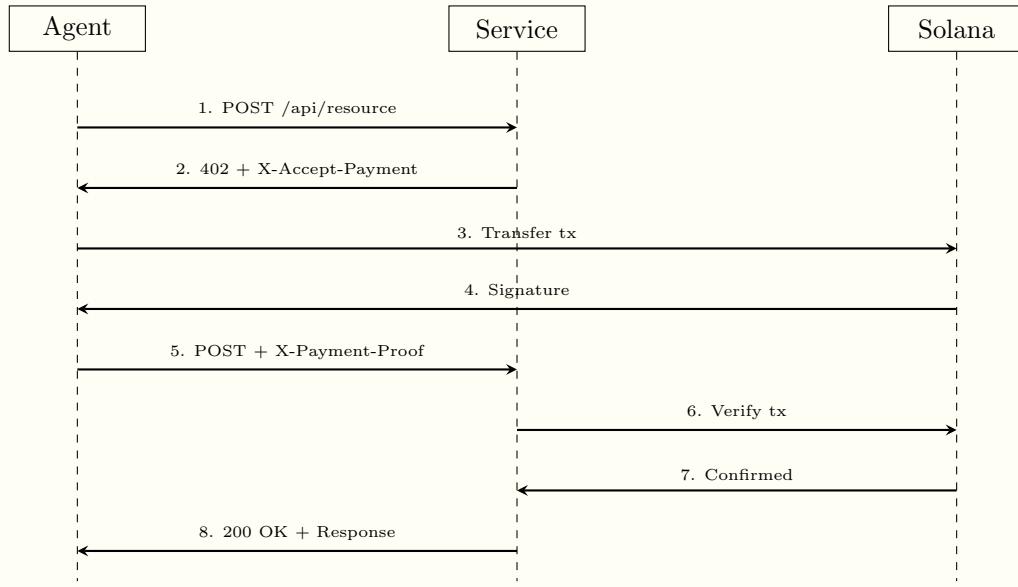


Figure 2: *x402 Payment Flow*. Agent receives 402, pays on-chain, then retries with proof.

4.3 Protocol Design

4.3.1 Payment Request

When a service requires payment, it responds with HTTP 402 and an X-Accept-Payment header:

```
1 HTTP/1.1 402 Payment Required
2 X-Accept-Payment: <base64-encoded-payment-request>
3 Content-Type: application/json
4
5 {
6     "error": "payment_required",
7     "message": "This endpoint requires payment"
8 }
```

Listing 3: HTTP 402 Payment Required Response

The payment request (base64-decoded) contains:

```
1 {
2     "version": "x0-v1",
3     "recipient": "7xKXtg2CW87d97TXJSDpbD5jBkheTqA83TZRuJosgAsU",
4     "amount": "1000000",
5     "resource": "/api/v1/generate",
6     "memo": "Text generation request",
```

```

7   "network": "solana-mainnet",
8   "escrow": {
9     "use_escrow": false,
10    "delivery_timeout": 3600,
11    "auto_release_delay": 86400,
12    "arbiter": null
13  }
14 }
```

Listing 4: Payment Request Structure

4.3.2 Payment Proof

After payment, the agent includes proof in subsequent requests:

```

1 POST /api/v1/generate HTTP/1.1
2 X-Payment-Proof: <base64-encoded-proof>
3 X-Payment-Version: x0-v1
4 Content-Type: application/json
5
6 {
7   "prompt": "Generate a haiku about blockchain"
8 }
```

Listing 5: HTTP Request with Payment Proof

Payment proof contains:

```

1 {
2   "signature": "5VDx8F...", // Transaction signature
3   "slot": 123456789,
4   "payer": "9xQeWv...",
5   "timestamp": 1706400000,
6   "network": "solana-mainnet"
7 }
```

Listing 6: Payment Proof Structure

4.3.3 Verification

The service verifies payment:

1. Fetch transaction by signature
2. Verify transaction succeeded
3. Check recipient matches service wallet
4. Check amount \geq requested amount
5. Verify timestamp within tolerance (± 5 minutes)
6. Check memo matches request (via SHA-256)

4.4 Security Properties

Theorem 4.1 (Payment Non-Repudiation). *A valid payment proof is unforgeable. An adversary cannot construct a proof without executing the on-chain transaction.*

Proof. The proof contains transaction signature $\sigma = \text{Sign}(sk_A, tx)$. By existential unforgeability of Ed25519, an adversary cannot produce σ' such that $\text{Verify}(pk_A, tx, \sigma') = 1$ without sk_A . On-chain verification ensures the transaction executed. \square \square

5 Conditional Escrow (x0-escrow)

5.1 Design

Escrow enables trustless payments for services with uncertain delivery.

5.1.1 Escrow State Machine

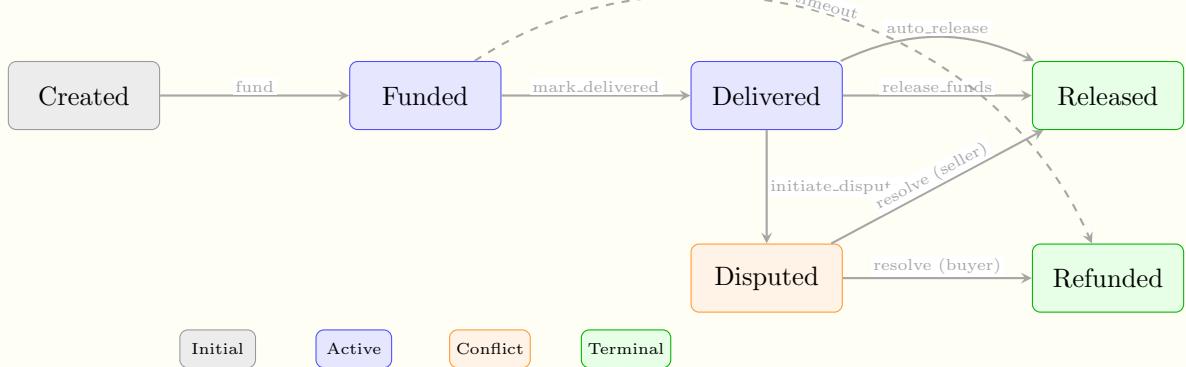


Figure 3: Escrow State Machine. States are colored by type: gray (initial), blue (active), orange (conflict), green (terminal).

5.1.2 Escrow Account

```

1 #[account]
2 pub struct EscrowAccount {
3     pub version: u8,                      // Account version
4     pub buyer: Pubkey,
5     pub seller: Pubkey,
6     pub arbiter: Option<Pubkey>,
7     pub amount: u64,
8     pub memo_hash: [u8; 32],
9     pub state: EscrowState,
10    pub timeout: i64,
11    pub created_at: i64,
12    pub delivery_proof: Option<[u8; 32]>,
13    pub dispute_evidence: Option<[u8; 32]>,
14    pub mint: Pubkey,
15    pub token_decimals: u8,
16    pub dispute_initiated_slot: u64, // MEDIUM-6: arbiter delay
17    pub bump: u8,

```

Listing 7: Escrow Account Structure

5.2 Key Operations

5.2.1 Create and Fund

Algorithm 2 Create and Fund Escrow

```

1: procedure CREATEESCROW( $B, S, A, x, m, \tau$ )
2:   require  $B \neq S$ 
3:   require  $3600 \leq \tau \leq 2592000$  ▷ 1h to 30d
4:    $e \leftarrow \text{EscrowPDA}(B, S, H(m))$ 
5:    $e.\text{timeout} \leftarrow t_{\text{now}} + \tau$ 
6:    $e.\text{state} \leftarrow \text{Created}$ 
7:   return  $e$ 
8: end procedure
9: procedure FUNDESCROW( $e, x$ )
10:  require  $e.\text{state} = \text{Created}$ 
11:  require  $t_{\text{now}} < e.\text{timeout}$ 
12:   $\text{transfer}(B, e, x)$ 
13:   $e.\text{state} \leftarrow \text{Funded}$ 
14: end procedure

```

5.2.2 Delivery and Release

Algorithm 3 Mark Delivered and Release Funds

```

1: procedure MARKDELIVERED( $e, p$ )
2:   require signer =  $e.\text{seller}$ 
3:   require  $e.\text{state} = \text{Funded}$ 
4:    $e.\text{delivery\_proof} \leftarrow H(p)$ 
5:    $e.\text{state} \leftarrow \text{Delivered}$ 
6: end procedure
7: procedure RELEASEFUNDS( $e$ )
8:   require signer =  $e.\text{buyer}$ 
9:   require  $e.\text{state} = \text{Delivered}$ 
10:   $e.\text{state} \leftarrow \text{Released}$  ▷ BEFORE transfer!
11:   $\text{transfer}(e, e.\text{seller}, e.\text{amount})$ 
12:   $\text{UpdateReputation}(e.\text{seller}, \text{Success})$  ▷ Optional CPI
13: end procedure

```

5.2.3 Auto-Release

After delivery, if buyer doesn't dispute within timeout, seller can claim:

Algorithm 4 Auto-Release After Timeout

```
1: procedure CLAIMAUTORELEASE( $e$ )
2:   require signer =  $e.\text{seller}$ 
3:   require  $e.\text{state} = \text{Delivered}$ 
4:   require  $t_{\text{now}} > e.\text{timeout}$ 
5:    $e.\text{state} \leftarrow \text{Released}$ 
6:   transfer( $e, e.\text{seller}, e.\text{amount}$ )
7:   UpdateReputation( $e.\text{seller}$ , Success)
8: end procedure
```

Algorithm 5 Initiate and Resolve Dispute

```
1: procedure INITIATEDISPUTE( $e, d$ )
2:   require signer  $\in \{e.\text{buyer}, e.\text{seller}\}$ 
3:   require  $e.\text{state} = \text{Delivered}$ 
4:    $e.\text{dispute\_evidence} \leftarrow H(d)$ 
5:    $e.\text{dispute\_initiated\_slot} \leftarrow \text{current\_slot}$ 
6:    $e.\text{state} \leftarrow \text{Disputed}$ 
7:   UpdateReputation( $e.\text{seller}$ , Dispute)
8: end procedure
9: procedure RESOLVEDISPUTE( $e, \text{release\_to\_seller}$ )
10:  require signer =  $e.\text{arbiter}$ 
11:  require  $e.\text{state} = \text{Disputed}$ 
12:  require  $\text{current\_slot} \geq e.\text{dispute\_initiated\_slot} + \Delta$ 
13:   $\triangleright \Delta = 216000 \text{ slots} \approx 24 \text{ hours}$ 
14:  if  $\text{release\_to\_seller}$  then
15:     $e.\text{state} \leftarrow \text{Released}$ 
16:    transfer( $e, e.\text{seller}, e.\text{amount}$ )
17:    UpdateReputation( $e.\text{seller}$ , ResolutionFavor)
18:  else
19:     $e.\text{state} \leftarrow \text{Refunded}$ 
20:    transfer( $e, e.\text{buyer}, e.\text{amount}$ )
21:  end if
22: end procedure
```

5.2.4 Dispute Resolution

5.3 Security Analysis

Theorem 5.1 (Escrow Safety). *An escrow satisfies:*

1. **Atomicity:** Funds are either fully released or refunded, never partially
2. **Liveness:** Funds are never permanently locked
3. **Fairness:** Either party can initiate dispute; arbiter is neutral

Proof. **Atomicity:** State transitions are atomic. The Released and Refunded states are terminal.

Liveness: Three exit paths exist:

- Buyer releases after delivery
- Seller claims auto-release after timeout
- Arbiter resolves dispute

At least one path is always available after delivery or timeout.

Fairness: Either party can initiate dispute. Arbiter resolution requires evidence from both parties and a 24-hour delay. □

6 Reputation Oracle (x0-reputation)

6.1 Design Goals

1. **Transparency:** All reputation data on-chain
2. **Temporal Decay:** Old successes shouldn't dominate forever
3. **Sybil Resistance:** New agents don't get perfect scores
4. **Nuanced Scoring:** Distinguish success, failure, disputes, resolutions

6.2 Reputation Account

```
1 #[account]
2 pub struct AgentReputation {
3     pub version: u8,                                // Account version (v2)
4     pub agent_id: Pubkey,
5     pub total_transactions: u64,
6     pub successful_transactions: u64,
7     pub disputed_transactions: u64,
8     pub resolved_in_favor: u64,
9     pub failed_transactions: u64,        // Policy rejections
10    pub average_response_time_ms: u32,
11    pub cumulative_response_time_ms: u64,
12    pub last_updated: i64,
13    pub last_decay_applied: i64,
14    pub bump: u8,
15 }
```

Listing 8: Reputation Account Structure

6.3 Reputation Score Calculation

Definition 6.1 (Reputation Score). The reputation score $S \in [0, 1]$ is computed as:

$$S = 0.60 \cdot S_{\text{success}} + 0.15 \cdot S_{\text{resolution}} \\ + 0.10 \cdot (1 - R_{\text{dispute}}) + 0.15 \cdot (1 - R_{\text{failure}}) \quad (11)$$

where:

$$S_{\text{success}} = \frac{n_{\text{success}}}{n_{\text{total}}} \quad (12)$$

$$S_{\text{resolution}} = \begin{cases} \frac{n_{\text{resolved}}}{n_{\text{disputed}}} & n_{\text{disputed}} > 0 \\ 0.5 & n_{\text{disputed}} = 0 \wedge n_{\text{total}} < 10 \\ 1.0 & n_{\text{disputed}} = 0 \wedge n_{\text{total}} \geq 10 \end{cases} \quad (13)$$

$$R_{\text{dispute}} = \frac{n_{\text{disputed}}}{n_{\text{total}}} \quad (14)$$

$$R_{\text{failure}} = \frac{n_{\text{failed}}}{n_{\text{total}}} \quad (15)$$

Remark 6.1. The neutral resolution score (0.5) for new agents prevents Sybil attacks where attackers create new identities to get perfect scores.

Example 6.1 (Reputation Score Calculation). Consider an agent with the following transaction history:

- Total transactions: $n_{\text{total}} = 100$
- Successful: $n_{\text{success}} = 90$
- Disputed: $n_{\text{disputed}} = 5$
- Resolved in favor: $n_{\text{resolved}} = 3$
- Failed (policy rejections): $n_{\text{failed}} = 2$

Computing component scores:

$$S_{\text{success}} = \frac{90}{100} = 0.90 \\ S_{\text{resolution}} = \frac{3}{5} = 0.60 \quad (\text{since } n_{\text{disputed}} > 0) \\ R_{\text{dispute}} = \frac{5}{100} = 0.05 \\ R_{\text{failure}} = \frac{2}{100} = 0.02$$

Final reputation score:

$$S = 0.60(0.90) + 0.15(0.60) + 0.10(1 - 0.05) + 0.15(1 - 0.02) \\ = 0.540 + 0.090 + 0.095 + 0.147 \\ = \mathbf{0.872}$$

This agent has a strong reputation (87.2%), with room for improvement in dispute resolution.

6.4 Temporal Decay

Definition 6.2 (Exponential Decay). Monthly decay applies to successful transactions:

$$n_{\text{success}}^{(t+1)} = n_{\text{success}}^{(t)} \cdot (1 - \alpha) \quad (16)$$

where $\alpha = 0.01$ (1% monthly decay) and t is measured in months.

For m months:

$$n_{\text{success}}^{(t+m)} = n_{\text{success}}^{(t)} \cdot (0.99)^m \quad (17)$$

Proposition 6.1 (Half-Life). *The half-life of reputation is:*

$$t_{1/2} = \frac{\ln 2}{\ln(1/0.99)} \approx 69 \text{ months} \quad (18)$$

Algorithm 6 Apply Reputation Decay

```

1: procedure APPLYDECAY( $r$ )
2:    $m \leftarrow \lfloor(t_{\text{now}} - r.\text{last\_decay\_applied})/2592000\rfloor$ 
3:   if  $m > 0$  then
4:      $d_{\text{mult}} \leftarrow 99^{\min(m, 12)}$                                  $\triangleright$  Cap at 12 months
5:      $d_{\text{div}} \leftarrow 100^{\min(m, 12)}$ 
6:      $r.\text{successful\_transactions} \leftarrow r.\text{successful\_transactions} \cdot d_{\text{mult}}/d_{\text{div}}$ 
7:      $r.\text{last\_decay\_applied} \leftarrow t_{\text{now}}$ 
8:   end if
9: end procedure

```

6.5 Authorization Model

Definition 6.3 (Authorized Callers). Reputation updates can only be called by:

- **Escrow program:** Records success/dispute/resolution
- **Guard program:** Records policy failures
- **Policy owner:** Self-reported off-chain transactions

Theorem 6.2 (Reputation Integrity). *An adversary cannot artificially inflate reputation without:*

1. *Completing escrow transactions (requires payment)*
2. *Winning disputes (requires arbiter approval)*
3. *Controlling policy owner key (equivalent to ownership)*

7 Agent Discovery (x0-registry)

7.1 Registry Entry

```

1 #[account]
2 pub struct AgentRegistry {
3     pub version: u8,                                // Account version
4     pub agent_id: Pubkey,                           // Policy PDA (primary ID)
5     pub owner: Pubkey,                            // Owner who can update
6     pub endpoint: String,                          // Service URL (max 256 chars)
7     pub capabilities: Vec<Capability>,           // Max 10 capabilities
8     pub price_oracle: Option<Pubkey>,            // Optional pricing oracle
9     pub reputation_pda: Pubkey,                  // Linked reputation account
10    pub last_updated: i64,
11    pub is_active: bool,
12    pub bump: u8,
13 }
14
15 pub struct Capability {
16     pub capability_type: String,                // Max 64 chars
17     pub metadata: String,                      // JSON metadata (max 256 chars)
18 }

```

Listing 9: Registry Entry Structure

The `Capability` struct uses a `JSON metadata` field to encode pricing, versioning, and service-specific parameters (see Section 7.3), avoiding rigid on-chain schema constraints.

7.2 Discovery Algorithm

Algorithm 7 Find Agents by Capability

```

1: procedure FINDAGENTS(cap_type)
2:   A ← {}
3:   for e ∈ GetProgramAccounts(registry_program) do
4:     if cap_type ∈ e.capabilities ∧ e.is_active then
5:       s ← FetchReputation(e.reputation_pda)
6:       A ← A ∪ {(e, s)}
7:     end if
8:   end for
9:   return SortByScore(A)
10: end procedure

```

7.3 Capability Metadata

Capabilities use JSON metadata:

```

1 {
2     "type": "text-generation",
3     "models": ["gpt-4", "claude-3"],
4     "pricing": {
5         "per_token": 0.00001,
6         "minimum": 0.01
7     },
8     "rate_limit": {
9         "requests_per_minute": 60,
10        "burst": 10
11    },

```

```

12     "api_version": "v1"
13 }
```

Listing 10: Capability Metadata Example

8 USDC Wrapper (x0-wrapper)

8.1 Problem Statement

Direct USDC usage has limitations:

1. No transfer hook support on standard USDC
2. Different token programs (SPL vs Token-2022)
3. Protocol fees require wrapper

8.2 Design

x0-USD is a 1:1 USDC-backed Token-2022 token with:

- Transfer hook pointing to x0-guard
- 0.8% transfer fee (configurable via PROTOCOL_FEE_BASIS_POINTS constant, default 80 bps)
- Configurable redemption fee (default 0.8%, set via WRAPPER_REDEMPTION_FEE_BPS)

8.2.1 State Accounts

The wrapper maintains two on-chain accounts for configuration and operational metrics:

```

1 #[account]
2 pub struct WrapperConfig {
3     pub admin: Pubkey,                                // Admin key (should be multisig)
4     pub pending_admin: Option<Pubkey>, // Two-step admin transfer
5     pub usdc_mint: Pubkey,                            // Underlying USDC mint
6     pub wrapper_mint: Pubkey,                         // x0-USD mint
7     pub reserve_account: Pubkey,                      // USDC reserve token account
8     pub redemption_fee_bps: u16,                      // Fee in basis points
9     pub is_paused: bool,                             // Emergency pause flag
10    pub bump: u8,
```

Listing 11: Wrapper Configuration Account

```

1 #[account]
2 pub struct WrapperStats {
3     pub reserve_usdc_balance: u64,      // Current USDC in reserve
4     pub outstanding_wrapper_supply: u64, // Outstanding x0-USD supply
5     pub total_deposits: u64,           // All-time deposits
6     pub total_redemptions: u64,         // All-time redemptions
7     pub total_fees_collected: u64,       // All-time fees
8     pub daily_redemption_volume: u64,   // Resets every 24h
9     pub daily_redemption_reset_timestamp: i64,
10    pub last_updated: i64,
```

```

11     pub bump: u8,
12 }
```

Listing 12: Wrapper Statistics Account

The reserve ratio ρ can be computed directly from `WrapperStats`:

$$\rho = \frac{\text{reserve_usdc_balance}}{\text{outstanding_wrapper_supply}} \quad (19)$$

8.2.2 Reserve Invariant

Definition 8.1 (Reserve Invariant). At all times, the following must hold:

$$R_{\text{USDC}} \geq S_{\text{x0-USD}} \quad (20)$$

where R_{USDC} is USDC reserve and $S_{\text{x0-USD}}$ is outstanding x0-USD supply.

Theorem 8.1 (Invariant Preservation). *The reserve invariant is preserved under all operations.*

Proof. We prove by cases:

Deposit: User deposits x USDC, receives x x0-USD.

$$R' = R + x, \quad S' = S + x \implies R' - S' = R - S \geq 0 \quad (21)$$

Redemption: User burns x x0-USD, receives $x - f$ USDC (fee f).

$$R' = R - (x - f), \quad S' = S - x \quad (22)$$

$$R' - S' = R - x + f - S + x = R - S + f \geq R - S \geq 0 \quad (23)$$

The fee increases the reserve ratio. □

8.3 Governance

8.3.1 Timelock

All admin operations require 48-hour timelock. Each pending action is stored in a PDA:

```

1 #[account]
2 pub struct AdminAction {
3     pub action_type: AdminActionType, // SetFeeRate | SetPaused |
4         EmergencyWithdraw | TransferAdmin
5     pub scheduled_timestamp: i64,      // Earliest execution time
6     pub new_value: u64,                // Interpretation depends on action_type
7     pub new_admin: Pubkey,            // For TransferAdmin actions
8     pub destination: Pubkey,          // For EmergencyWithdraw
9     pub executed: bool,
10    pub cancelled: bool,
11    pub bump: u8,
```

Listing 13: Timelocked Admin Action

An action is executable if and only if $t_{\text{now}} \geq \text{scheduled_timestamp}$ and $\neg \text{executed} \wedge \neg \text{cancelled}$.

Algorithm 8 Timelock Pattern

```
1: procedure SCHEDULEACTION( $a, v, t$ )
2:   action_pda  $\leftarrow$  CreateActionPDA( $a, \text{nonce}$ )
3:   action_pda.type  $\leftarrow a$ 
4:   action_pda.value  $\leftarrow v$ 
5:   action_pda.scheduled_time  $\leftarrow t + 172800$  ▷ 48h
6: end procedure
7: procedure EXECUTEACTION(action_pda)
8:   require  $t_{\text{now}} \geq \text{action\_pda}.scheduled\_time$ 
9:   require  $\neg \text{action\_pda}.executed$ 
10:  require  $\neg \text{action\_pda}.cancelled$ 
11:  ApplyAction(action_pda.type, action_pda.value)
12:  action_pda.executed  $\leftarrow \text{true}$ 
13: end procedure
```

8.3.2 Emergency Pause

Emergency pause is the **only** operation bypassing timelock:

- Can only **pause**, never unpause
- Unpausing requires standard timelock
- Prevents rapid pause/unpause attacks

8.4 Reserve Ratio Monitoring

Definition 8.2 (Reserve Ratio).

$$\rho = \frac{R_{\text{USDC}}}{S_{\text{x0-USD}}} \quad (24)$$

Alert levels:

- $\rho < 1.01$: Warning alert
- $\rho < 1.00$: Critical alert (undercollateralized)

9 Human-in-the-Loop (Blinks)

9.1 Problem Statement

Full automation is dangerous for:

1. Transfers exceeding daily limit
2. Recipients not on whitelist
3. Unusually large transactions

9.2 Blink Design

A **Blink** is a Solana Action requesting human approval.

9.2.1 Generation

When guard rejects a transfer due to limits, it can:

1. Generate Blink ID: $b = H(\text{policy} \parallel R \parallel x \parallel t)[0..16]$
2. Emit `BlinkGenerated` event with:
 - Blink ID
 - Requested amount
 - Recipient
 - Expiration (15 minutes)
3. Charge fee (0.001 SOL) to prevent spam

9.2.2 Rate Limiting

Definition 9.1 (Blink Rate Limit). Maximum 3 Blanks per hour per policy.

Algorithm 9 Blink Rate Limit Check

```
1: procedure CHECKBLINKRATELIMIT( $p, t$ )
2:    $h_{\text{current}} \leftarrow \lfloor t/3600 \rfloor$ 
3:    $h_{\text{window}} \leftarrow \lfloor p.\text{blink\_hour\_start}/3600 \rfloor$ 
4:   if  $h_{\text{current}} \neq h_{\text{window}}$  then
5:      $p.\text{blink\_hour\_start} \leftarrow t$ 
6:      $p.\text{blinks\_this\_hour} \leftarrow 0$ 
7:   end if
8:   if  $p.\text{blinks\_this\_hour} \geq 3$  then
9:     return false
10:  end if
11:   $p.\text{blinks\_this\_hour} \leftarrow p.\text{blinks\_this\_hour} + 1$ 
12:  return true
13: end procedure
```

9.2.3 Approval

Owner approves via:

1. Viewing Blink details (QR code, web interface, wallet)
2. Signing approval transaction
3. Temporarily raising limit or whitelisting recipient

10 Economic Model

10.1 Fee Structure

10.2 Fee Collection

Transfer fees are collected via Token-2022's TransferFee extension:

Operation	Fee	Recipient
Transfer (x0-USD)	0.8%	Protocol treasury
Registry listing	0.1 SOL	Protocol treasury
Blink generation	0.001 SOL	Protocol treasury
Wrapper redemption	0.8% (configurable)	Wrapper reserve

Table 1: Protocol Fee Schedule

Algorithm 10 Fee Harvesting

```

1: procedure HARVESTFEES(mint, accounts)
2:   for  $a \in \text{accounts}$  do
3:      $f \leftarrow \text{GetWithheldAmount}(a)$ 
4:      $\text{TransferWithheldToMint}(a, f)$ 
5:   end for
6:    $F \leftarrow \sum_a f$ 
7:    $\text{WithdrawWithheldFromMint}(\text{mint}, \text{treasury}, F)$ 
8: end procedure

```

10.3 Token Economics

10.3.1 x0-USD Supply Dynamics

$$\frac{dS}{dt} = D(t) - R(t) \quad (25)$$

$$\frac{dR_{\text{USDC}}}{dt} = D(t) - R(t) + f \cdot R(t) \quad (26)$$

where:

- $S(t)$ = x0-USD supply
- $R_{\text{USDC}}(t)$ = USDC reserve
- $D(t)$ = Deposit rate
- $R(t)$ = Redemption rate (before fees)
- f = Redemption fee rate

Reserve ratio evolution:

$$\rho(t) = \frac{R_{\text{USDC}}(t)}{S(t)} = 1 + \int_0^t \frac{f \cdot R(\tau)}{S(\tau)} d\tau \quad (27)$$

The reserve ratio increases over time due to fees.

11 Zero-Knowledge Proof Verification (x0-zk-verifier)

11.1 Problem Statement

Token-2022's Confidential Transfer extension encrypts token balances and transfer amounts using twisted ElGamal encryption over the Ristretto255 curve. To ensure that encrypted operations preserve token invariants (e.g., non-negative balances, valid public keys), the protocol

requires zero-knowledge proofs that attest to the correctness of ciphertext operations without revealing plaintext values.

The challenge is two-fold:

1. **Client-side proof generation:** Proofs must be generated off-chain where the user holds the ElGamal secret key, using Groth16 proving circuits from `solana-zk-token-sdk`.
2. **On-chain proof verification:** Proofs must be verified on-chain and the verification result stored in a state account that Token-2022 can reference during confidential transfer execution.

11.2 Cryptographic Foundations

11.2.1 Twisted ElGamal Encryption

Token-2022 uses a *twisted* variant of ElGamal encryption over the Ristretto255 group, which provides both encryption and efficient homomorphic addition:

Definition 11.1 (Twisted ElGamal Encryption). For a public key $P = s \cdot G$ where s is the secret scalar and G is the Ristretto255 basepoint:

$$\text{Encrypt}(P, v) = (r \cdot G, r \cdot P + v \cdot H) \quad (28)$$

$$\text{Decrypt}(s, (C_1, C_2)) = C_2 - s \cdot C_1 = v \cdot H \quad (29)$$

where $r \xleftarrow{\$} \mathbb{Z}_q$ is a random scalar, H is a second generator, and v is the plaintext value. Decryption recovers $v \cdot H$, from which v is obtained via brute-force discrete log (feasible for $v < 2^{48}$).

11.2.2 Groth16 Proof System

The protocol uses Groth16 [11] for succinct non-interactive zero-knowledge proofs:

Definition 11.2 (Groth16 Proof). A Groth16 proof $\pi = (A, B, C) \in \mathbb{G}_1 \times \mathbb{G}_2 \times \mathbb{G}_1$ proves knowledge of a witness w satisfying a rank-1 constraint system $R(x, w) = 1$, where x is the public input. The proof satisfies:

- **Completeness:** An honest prover always produces a valid proof
- **Soundness:** No efficient adversary can produce a valid proof for a false statement
- **Zero-knowledge:** The proof reveals nothing about w beyond $R(x, w) = 1$

Proof size is constant: 192 bytes regardless of circuit complexity.

11.3 Proof Types

The verifier supports four proof types corresponding to Token-2022 confidential transfer operations:

11.3.1 PubkeyValidityProof

Proves that a given 32-byte value is a valid ElGamal public key (i.e., a point on the Ristretto255 curve that was correctly derived from a secret key).

$$\text{Prove} : \exists s \in \mathbb{Z}_q \text{ such that } P = s \cdot G \quad (30)$$

Required when configuring a token account for confidential transfers. Proof data is 64 bytes.

11.3.2 ZeroBalanceProof

Proves that an ElGamal ciphertext encrypts the value zero, required for account closure:

$$\text{Prove} : (C_1, C_2) = \text{Encrypt}(P, 0) \implies C_2 = r \cdot P \quad (31)$$

Proof data is 96 bytes.

11.3.3 WithdrawProof

Proves that a withdrawal of amount a from an encrypted balance B is valid:

$$\text{Prove} : \text{Decrypt}(s, B) = v \geq a \wedge B' = B - \text{Encrypt}(P, a) \quad (32)$$

where v is the current balance and B' is the new encrypted balance. This ensures the account has sufficient funds without revealing either the balance or the withdrawal amount on-chain. Proof data is 160 bytes.

11.3.4 TransferProof

Proves the correctness of a confidential transfer between two accounts:

$$\text{Prove} : \text{Decrypt}(s_{\text{src}}, B_{\text{src}}) \geq a \wedge B'_{\text{src}} = B_{\text{src}} - \text{Encrypt}(P_{\text{src}}, a) \wedge \Delta_{\text{dst}} = \text{Encrypt}(P_{\text{dst}}, a) \quad (33)$$

This is the most complex proof, attesting simultaneously that the sender has sufficient balance, the sender's balance is correctly decremented, and the recipient's balance increment is correctly encrypted under the recipient's public key.

11.4 Architecture

11.4.1 Proof Context Account

Successful verification creates a `ProofContext` PDA that stores the verification result:

```
1 #[account]
2 pub struct ProofContext {
3     pub version: u8,
4     pub proof_type: ProofType,           // PubkeyValidity | Withdraw |
5     pub verified: bool,
6     pub owner: Pubkey,                 // Account owner who can use this proof
7     pub verified_at: i64,               // Unix timestamp
```

```

8  pub amount: Option<u64>,           // For Withdraw/Transfer proofs
9  pub recipient: Option<Pubkey>,      // For Transfer proofs
10 pub elgamal_pubkey: Option<[u8; 32]>, // For PubkeyValidity proofs
11 pub mint: Pubkey,                  // Token mint this proof is for
12 pub token_account: Pubkey,        // Token account this proof is for
13 pub bump: u8,
14 }

```

Listing 14: Proof Context State Account

11.4.2 Proof Freshness

To prevent replay attacks, proof contexts enforce a temporal validity window:

Definition 11.3 (Proof Freshness). A proof context C is *fresh* at time t if and only if:

$$t - C.\text{verified_at} < \Delta_{\text{proof}} = 300 \text{ seconds} \quad (34)$$

The 5-minute window provides sufficient time for transaction construction and submission while limiting the replay window.

11.4.3 Amount Bounds

All confidential amounts are bounded:

$$0 \leq a \leq 2^{48} - 1 = 281,474,976,710,655 \quad (35)$$

This constraint arises from the brute-force discrete log required during ElGamal decryption. The `solana-zk-token-sdk` uses a precomputed lookup table for values up to 2^{16} , combined with baby-step-giant-step for the remaining 2^{32} range.

11.5 Off-Chain Proof Generation (x0-zk-proofs)

The `x0-zk-proofs` crate is a Rust library compiled to WebAssembly via `wasm-pack`, providing client-side proof generation for the TypeScript SDK:

```

1 #[wasm_bindgen]
2 pub fn generate_pubkey_validity_proof(
3     keypair_bytes: &[u8]           // 64-byte ElGamalKeypair
4 ) -> Result<Vec<u8>, JsValue>; // 64-byte PubkeyValidityData
5
6 #[wasm_bindgen]
7 pub fn generate_zero_balance_proof(
8     keypair_bytes: &[u8]           // 64-byte ElGamalKeypair
9 ) -> Result<Vec<u8>, JsValue>; // 96-byte ZeroBalanceProofData
10
11 #[wasm_bindgen]
12 pub fn generate_withdraw_proof(
13     keypair_bytes: &[u8],         // 64-byte ElGamalKeypair
14     current_balance: u64,        // Current decrypted balance
15     withdraw_amount: u64         // Amount to withdraw
16 ) -> Result<Vec<u8>, JsValue>; // 160-byte WithdrawData

```

Listing 15: x0-zk-proofs WASM Interface

The WASM module internally uses:

- `solana-zk-token-sdk` v1.18.22 for Groth16 circuit construction and proving
- `ElGamalKeypair` reconstruction from 64-byte serialized form
- `bytemuck` for zero-copy POD serialization of proof data
- Authenticated encryption (AE) ciphertexts for decryptable balance hints

11.6 Verification Flow

The end-to-end flow for a confidential withdrawal is:

Algorithm 11 Confidential Withdrawal with ZK Proof

```

1: procedure CONFIDENTIALWITHDRAW(keypair,  $b$ ,  $a$ )
2:   require  $a \leq b$                                  $\triangleright$  Client-side: sufficient balance
3:   require  $a \leq 2^{48} - 1$                        $\triangleright$  Client-side: amount bound
4:    $\pi \leftarrow \text{x0_zk_proofs::generate_withdraw_proof}(\text{keypair}, b, a)$ 
5:   send x0_zk_verifier::verify_withdraw( $\pi, a, \text{new\_ae\_balance}$ )
6:    $C \leftarrow \text{ProofContext PDA}$                    $\triangleright$  On-chain: created after verification
7:   require  $C.\text{verified} = \text{true}$              $\triangleright$  On-chain check
8:   send Token-2022::confidential_withdraw( $a, C$ )
9: end procedure

```

11.7 Security Properties

Theorem 11.1 (Balance Confidentiality). *An adversary observing the blockchain cannot determine any agent's confidential balance, even given all historic proof contexts and ciphertexts.*

Proof. Token balances are encrypted under twisted ElGamal with semantic security. Proof contexts store only the verification result, proof type, and amount (for withdrawals), but the *remaining balance* is revealed only via AE ciphertext decryptable solely by the account holder. The Groth16 zero-knowledge property ensures proofs reveal nothing beyond the statement's truth. \square \square

Theorem 11.2 (Proof Unforgeability). *No computationally bounded adversary can produce a valid proof for a false statement (e.g., withdrawing more than the encrypted balance).*

Proof. Follows from the knowledge-soundness of Groth16 under the q -type knowledge-of-exponent assumption in bilinear groups. The verifier re-derives the check via `solana-zk-token-sdk`'s `verify()`, which evaluates the Groth16 pairing equation. \square \square

12 Security Analysis

12.1 Threat Model

We consider the following adversaries:

1. **Compromised Agent:** Attacker obtains sk_A

2. **Malicious Service:** Service provider attempts fraud
3. **Malicious Agent:** Agent attempts reputation manipulation
4. **Wrapper Attack:** Attacker attempts reserve drain
5. **Bridge Attack:** Attacker attempts cross-chain fund extraction (see Section 14.10 for bridge-specific analysis)

12.2 Attack Scenarios and Mitigations

12.2.1 Compromised Agent Key

Attack: Adversary steals sk_A , attempts unlimited spending.

Mitigation:

- Rolling window limits spending to L per 24 hours
- Owner can revoke via `revoke_agent_authority`
- Maximum loss bounded by $L + x_{\max}$ where x_{\max} is max transaction size

12.2.2 Sybil Attack on Reputation

Attack: Adversary creates multiple identities to appear trustworthy.

Mitigation:

- Registry listing costs 0.1 SOL
- New agents get neutral (0.5) resolution score, not perfect
- Temporal decay prevents dormant identities

Cost for 100 Sybil identities: $100 \times 0.1 = 10 \text{ SOL} \approx \1000 (at \$100/SOL).

12.2.3 Escrow Griefing

Attack: Malicious buyer creates many escrows, ties up seller funds.

Mitigation:

- Escrow creation is permissionless but buyer must fund
- Sellers can set minimum transaction amounts
- Reputation tracking penalizes frivolous disputes

12.2.4 Wrapper Reserve Drain

Attack: Exploit bug to withdraw USDC without burning x0-USD.

Mitigation:

- Reserve invariant checked before every redemption
- State updated before transfer (reentrancy protection)
- Admin operations timelocked (48h notice)
- Emergency pause available

12.2.5 Transfer Hook Configuration Attack

Attack: Unauthorized party front-runs transfer hook initialization for a mint, potentially misconfiguring the extra account metas or consuming the PDA address.

Mitigation:

- Only the Token-2022 mint authority can initialize extra account metas
- Re-initialization is blocked after first valid configuration
- Extra account metas configuration is deterministic (not caller-supplied)
- Mint authority verification uses `StateWithExtensions` for robust parsing

Impact: Even though the extra metas configuration is hardcoded in the program, preventing unauthorized initialization eliminates a griefing vector where an attacker could pay the rent and claim the PDA before the legitimate mint authority initializes it.

12.3 Formal Verification

12.3.1 Spend Limit Invariant

Theorem 12.1 (Daily Spend Limit). *For all execution traces, if policy has daily limit L :*

$$\forall t : \sum_{i:t_i > t - 86400} x_i \leq L \quad (36)$$

Proof. Proven in Section 3.4. □

12.3.2 Reserve Invariant

Theorem 12.2 (Wrapper Reserve Invariant). *For all execution traces:*

$$\forall t : R_{USDC}(t) \geq S_{x0\text{-USD}}(t) \quad (37)$$

Proof. Proven in Section 7.2. □

13 Performance Analysis

13.1 Computational Complexity

13.2 On-Chain Costs

13.3 Transaction Benchmarks

Measured on Solana devnet (February 2026):

14 Related Work

14.1 Payment Channels

Bitcoin’s Lightning Network [4] and Ethereum’s Raiden [5] enable off-chain micropayments. Limitations:

Operation	Time	Space
Transfer validation	$O(n)$	$O(n)$
Merkle verification	$O(\log m)$	$O(\log m)$
Bloom verification	$O(k)$	$O(1)$
Reputation update	$O(1)$	$O(1)$
Registry search	$O(N)$	$O(1)$
STARK proof verification	$O(1)$	$O(\pi)$
Event log validation	$O(E \cdot T)$	$O(1)$

Table 2: Computational Complexity (n = window size, m = whitelist size, k = hash count, N = registered agents, E = event count, T = topics per event, $|\pi|$ = proof size)

Account	Size (bytes)	Rent (SOL)
AgentPolicy	4,952	0.035
EscrowAccount	307	0.0030
AgentRegistry	2,064	0.015
AgentReputation	114	0.0016
WrapperConfig	243	0.0025
WrapperStats	137	0.0018
ProofContext	248	0.0026
AdminAction	120	0.0016
BridgeConfig	589	0.0053
BridgeMessage	200	0.0022
EVMProofContext	3,180	0.024
BridgeOutMessage	153	0.0019
BridgeAdminAction	148	0.0018

Table 3: Account Sizes and Rent Costs (at 0.00000348 SOL/byte-epoch)

Operation	Compute Units	Latency
Initialize policy	45,000	450ms
Validate transfer (Merkle)	32,000	320ms
Validate transfer (Bloom)	18,000	180ms
Create escrow	28,000	280ms
Registry listing	35,000	350ms
Reputation update	15,000	150ms
Handle message (bridge)	42,000	420ms
Verify EVM proof (STARK)	500,000	5,000ms
Execute mint (bridge)	200,000	2,000ms
Initiate bridge out	180,000	1,800ms

Table 4: Transaction Performance (1 CU \approx 10 μ s)

- Require channel opening/closing (on-chain)
- Liquidity fragmentation
- Limited programmability

x0 uses on-chain transactions but adds programmable policy enforcement unavailable in payment channels.

14.2 Escrow Protocols

Traditional escrow (e.g., Escrow.com) requires trusted third parties. Smart contract escrow (OpenBazaar, LocalBitcoins) removes intermediaries but lacks:

- Reputation integration
- Automated dispute resolution
- Cross-protocol standards

14.3 Reputation Systems

OpenBazaar and Augur implement on-chain reputation but lack:

- Temporal decay
- Sybil resistance for new participants
- Integration with payment protocols

x0's reputation oracle addresses these limitations.

15 FROSTGATE: Trustless Cross-Chain Bridge

Remark 15.1 (Design Scope). FROSTGATE is *not* a mandatory component of the x0 protocol. The core x0 system (Sections 3–13) operates independently on Solana, providing programmable spending policies, escrow, reputation, and USDC wrapping without any cross-chain dependencies. FROSTGATE is presented as a theoretical extension for trustless interchain agentic interactions, demonstrating how x0's architecture can be extended to support Base \longleftrightarrow Solana payments while preserving the protocol's security guarantees. Deployments may choose to operate x0 purely on Solana, or integrate FROSTGATE to enable cross-chain agent commerce.

15.1 Problem Statement

Autonomous agent payments are inherently multi-chain. An agent purchasing compute on Base must pay a service provider on Solana, and vice versa. Existing bridge architectures—relying on multisig committees (Wormhole), optimistic fraud windows (Across), or centralized attestors (LayerZero)—introduce trust assumptions incompatible with the permissionless, real-time settlement guarantees agents require. A compromised bridge committee can fabricate deposits, drain reserves, and halt all cross-chain commerce.

We introduce *FROSTGATE*, a bidirectional Base \longleftrightarrow Solana bridge that eliminates all off-chain trust assumptions by combining:

1. **Hyperlane message passing** for cross-chain message delivery with on-chain sender verification,
2. **SP1 STARK proofs** for trustless verification of source-chain state transitions, and
3. **CPI-enforced reserve invariants** through the existing x0-wrapper mint/burn interface (Section 8).

The result is a bridge where every cross-chain transfer is backed by a succinct cryptographic proof of the source-chain event, verified *on-chain* by the destination chain, with no trusted relayer, committee, or oracle in the critical path.

15.2 Definitions

Definition 15.1 (Cross-Chain Bridge). A cross-chain bridge $\mathcal{B} = (\text{Lock}, \text{Relay}, \text{Prove}, \text{Mint})$ is a tuple of protocols that transfers value from a source chain \mathcal{C}_s to a destination chain \mathcal{C}_d such that for every deposit of amount a on \mathcal{C}_s , exactly a units of the wrapped asset are minted on \mathcal{C}_d , and no minting occurs without a corresponding valid deposit.

Definition 15.2 (Bridge Soundness). A bridge \mathcal{B} is *sound* if for all probabilistic polynomial-time adversaries \mathcal{A} :

$$\Pr[\text{Mint}(a', r') = 1 \mid \# \text{Lock}(a', r') \text{ on } \mathcal{C}_s] \leq \text{negl}(\lambda) \quad (38)$$

where a' is the minted amount, r' is the recipient, and λ is the security parameter.

Definition 15.3 (Bridge Completeness). A bridge \mathcal{B} is *complete* if for every valid deposit $\text{Lock}(a, r)$ on \mathcal{C}_s that is finalized, there exists a protocol execution that produces $\text{Mint}(a, r)$ on \mathcal{C}_d .

15.3 Architecture Overview

FROSTGATE operates as two mirror-image protocols:

- **Inbound** (Base → Solana): Lock USDC on Base, prove the lock via SP1 STARK, mint x0-USD on Solana.
- **Outbound** (Solana → Base): Burn x0-USD on Solana, prove the burn via SP1 STARK, unlock USDC on Base.

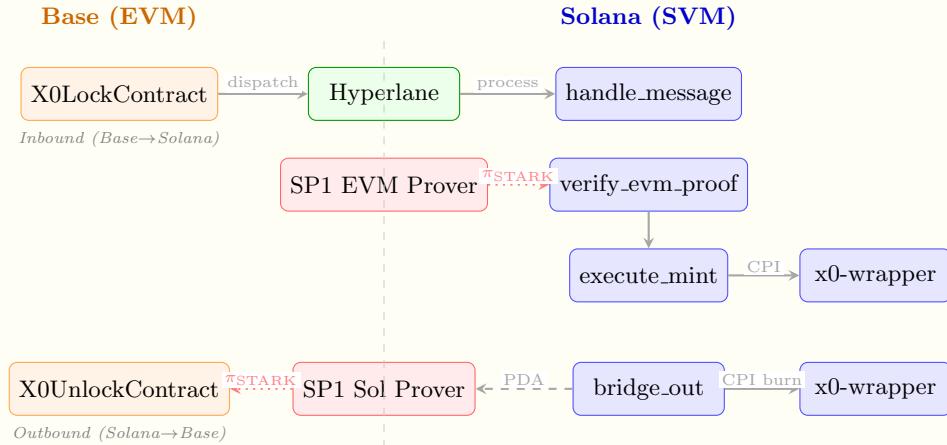


Figure 2: FROSTGATE Bridge Architecture. Orange nodes are EVM contracts. Blue nodes are Solana programs. Red nodes are off-chain SP1 provers.

15.4 On-Chain State

15.4.1 Bridge Configuration

The bridge is governed by a singleton PDA seeded by ["bridge_config"]:

```

1 pub struct BridgeConfig {
2     pub admin: Pubkey,           // Multisig administrator
3     pub hyperlane_mailbox: Pubkey, // Hyperlane mailbox program
4     pub sp1_verifier: Pubkey,    // SP1 verifier program
5     pub wrapper_program: Pubkey, // x0-wrapper for CPI
6     pub wrapper_config: Pubkey, // x0-wrapper config PDA
7     pub usdc_mint: Pubkey,      // USDC mint on Solana
8     pub wrapper_mint: Pubkey,   // x0-USD mint
9     pub bridge_usdc_reserve: Pubkey, // USDC reserve token account
10    pub is_paused: bool,
11    pub total_bridged_in: u64,    // All-time inflow (USDC micro-units)
12    pub total_bridged_out: u64,   // All-time outflow
13    pub nonce: u64,              // Monotonic message nonce
14    pub daily_inflow_volume: u64, // Rolling 24h inflow
15    pub daily_outflow_volume: u64, // Rolling 24h outflow
16    pub allowed_evm_contracts: [[u8; 20]; 10], // Whitelisted EVM contracts
17    pub supported_domains: [u32; 10],           // Hyperlane domain IDs
18    pub bridge_out_nonce: u64,      // Outbound monotonic nonce
19 }

```

Listing 16: BridgeConfig State Account

15.4.2 Bridge Message

Each inbound Hyperlane message creates a PDA seeded by ["bridge_message", message_id]:

```

1 pub struct BridgeMessage {
2     pub message_id: [u8; 32],    // SHA-256(origin || sender || body)
3     pub origin_domain: u32,      // e.g., 8453 (Base)
4     pub sender: [u8; 32],        // EVM address zero-padded to 32 bytes
5     pub recipient: Pubkey,       // Solana destination
6     pub amount: u64,             // USDC micro-units (6 decimals)
7     pub status: BridgeMessageStatus, // Received -> ProofVerified -> Minted
8     pub evm_tx_hash: [u8; 32],   // Source transaction hash
9     pub nonce: u64,              // Sender nonce for ordering
10 }

```

Listing 17: BridgeMessage State Account

The message status follows a linear state machine:

$$\text{Received} \xrightarrow{\text{verify_evm_proof}} \text{ProofVerified} \xrightarrow{\text{execute_mint}} \text{Minted} \quad (39)$$

Each transition is irreversible and guarded by status preconditions, ensuring no double-minting.

15.4.3 EVM Proof Context

Successful STARK verification creates a PDA seeded by ["evm_proof", message_id]:

```

1 pub struct EVMPProofContext {
2     pub proof_type: EVMPProofType,
3     pub verified: bool,
4     pub verified_at: i64,           // Clock timestamp
5     pub block_hash: [u8; 32],      // Proven EVM block hash
6     pub block_number: u64,
7     pub tx_hash: [u8; 32],         // Proven transaction hash
8     pub from: [u8; 20],           // EVM sender
9     pub to: [u8; 20],             // EVM contract address
10    pub value: u64,
11    pub event_logs: Vec<EVMEEventLog>, // Extracted receipt logs
12    pub message_id: [u8; 32],       // Links to BridgeMessage
13 }

```

Listing 18: EVMPProofContext State Account

15.4.4 Bridge Out Message

Each outbound transfer creates a PDA seeded by ["bridge_out_message", nonce.to_le_bytes()]:

```

1 pub struct BridgeOutMessage {
2     pub nonce: u64,                // Monotonic from BridgeConfig
3     pub solana_sender: Pubkey,     // Burn initiator
4     pub evm_recipient: [u8; 20],   // Base destination address
5     pub amount: u64,               // USDC micro-units
6     pub burn_tx_signature: [u8; 32], // Solana transaction signature
7     pub burned_at: i64,            // Burn timestamp
8     pub status: BridgeOutStatus,  // Burned -> Unlocked
9 }

```

Listing 19: BridgeOutMessage State Account

15.5 Inbound Protocol (Base → Solana)

15.5.1 Step 1: USDC Lock on Base

The XOLockContract (Solidity, OpenZeppelin) receives USDC via `transferFrom`, stores it in the contract, and dispatches a Hyperlane message:

Definition 15.4 (Bridge Message Body). The message body is a tightly-packed 80-byte encoding:

$$\text{body} = \text{solanaRecipient}_{[32]} \parallel \text{amount}_{\text{BE}[8]} \parallel \text{txHash}_{[32]} \parallel \text{nonce}_{\text{BE}[8]} \quad (40)$$

where \parallel denotes concatenation and subscripts indicate byte lengths. Big-endian encoding is used for all integer fields to match EVM conventions.

The lock emits a `Locked` event with indexed parameters for efficient filtering:

$$\text{Locked}(\underbrace{\text{sender}}_{\text{indexed}}, \underbrace{\text{solanaRecipient}}_{\text{indexed}}, \text{amount}, \text{nonce}, \text{messageId}) \quad (41)$$

The event signature is the Keccak-256 hash of the full event declaration:

$$\text{topics}[0] = \text{keccak256}("Locked(address,bytes32,uint256,uint256,bytes32)") \quad (42)$$

This hash is stored as a protocol constant and used by both the SP1 circuit and the on-chain verifier to locate the deposit event in receipt logs.

15.5.2 Step 2: Hyperlane Message Delivery

The Hyperlane protocol relays the message body to Solana, invoking the bridge program's `handle_message` instruction. The instruction validates:

Algorithm 12 Inbound Message Reception (`handle_message`)

```

1: procedure HANDLEMESSAGE(origin, sender, body)
2:   require  $\neg$  config.is_paused                                 $\triangleright$  Bridge not paused
3:   require ValidateProcessAuthority(caller)            $\triangleright$  Hyperlane PDA
4:   require origin  $\in$  config.supported_domains       $\triangleright$  Known domain
5:   require sender[0..12] =  $0^{12}$                        $\triangleright$  EVM format
6:   evm_addr  $\leftarrow$  sender[12..32]
7:   require evm_addr  $\in$  config.allowed_evm_contracts     $\triangleright$  Whitelist
8:   require  $|body| = 80$                                       $\triangleright$  Exact body size
9:   (recipient,  $a$ , tx_hash,  $n$ )  $\leftarrow$  decode(body)
10:  require  $a \in [\text{MIN\_AMOUNT}, \text{MAX\_AMOUNT}]$            $\triangleright$  Bounds check
11:  require daily_inflow +  $a \leq \text{MAX\_DAILY\_INFLOW}$          $\triangleright$  Rate limit
12:  id  $\leftarrow$  SHA256(origin  $\parallel$  sender  $\parallel$  body)
13:  create BridgeMessage(id, origin, sender, recipient,  $a$ )
14: end procedure

```

Remark 15.2 (Hyperlane Process Authority). The bridge validates the caller by deriving the Hyperlane process authority PDA:

$$\text{process_authority} = \text{PDA}_{\text{mailbox}}(\text{"hyperlane"} \parallel \text{"-"} \parallel \text{"process_authority"} \parallel \text{"-"} \parallel \text{bridge_id}) \quad (43)$$

This ensures only messages routed through the Hyperlane mailbox can trigger bridge message creation, even if an adversary constructs a valid-looking message body.

15.5.3 Step 3: STARK Proof Verification

After the BridgeMessage is created, an off-chain relayer generates an SP1 STARK proof of the EVM transaction and submits it via `verify_evm_proof`. This instruction:

1. Invokes the SP1 verifier program via CPI to verify the STARK proof
2. Decodes the public inputs: $(h_{\text{block}}, n_{\text{block}}, h_{\text{tx}}, from, to, v, success, logs)$
3. Validates `success = true` (EVM transaction succeeded)
4. Validates `tx_hash = BridgeMessage.evm_tx_hash`
5. **Validates the Locked event via `validate_locked_event()`**
6. Creates the `EVMProofContext` PDA and transitions the message to `ProofVerified`

Definition 15.5 (Locked Event Validation). Given a set of proven event logs \mathcal{L} and a BridgeMessage M , the function `validate_locked_event(\mathcal{L}, M)` succeeds if and only if there exists $\ell \in \mathcal{L}$ such that all of the following hold:

1. $\ell.contract \in \text{config.allowed_evm_contracts}$ (trusted contract)
2. $\ell.topics[0] = \text{LOCKED_EVENT_SIGNATURE}$ (correct event type)
3. $|\ell.topics| \geq 3$ (has indexed params)
4. $\ell.topics[2] = M.recipient$ (recipient match)
5. $|\ell.data| \geq 96$ (three ABI words)
6. $\text{abi.decode}(\ell.data).amount = M.amount$ (amount match)
7. $\text{abi.decode}(\ell.data).nonce = M.nonce$ (nonce match)

Theorem 15.1 (Event Validation Soundness). *Under the collision resistance of Keccak-256 and the soundness of the SP1 STARK proof system, no computationally bounded adversary can cause `validate_locked_event` to accept for a BridgeMessage M unless a genuine **Locked** event with matching recipient, amount, and nonce was emitted by a whitelisted contract in a finalized Base transaction.*

Proof. Suppose an adversary \mathcal{A} causes acceptance without a genuine event. Then either:

1. \mathcal{A} forged the SP1 proof (contradicts STARK soundness with $\leq 2^{-100}$ forgery probability),
2. \mathcal{A} found a collision in the event signature (contradicts Keccak-256 collision resistance), or
3. \mathcal{A} caused a whitelisted contract to emit a **Locked** event with the adversary's chosen parameters—but this requires calling `lock()` with matching `amount` and `solanaRecipient`, which requires the adversary to actually deposit USDC, making the event genuine.

All cases lead to contradiction. □

15.5.4 Step 4: Mint Execution

Once the proof is verified, `execute_mint` performs the mint via two CPIs:

Remark 15.3 (Wrapper Reserve Invariant Preservation). The two-step CPI sequence preserves the wrapper reserve invariant (Theorem 8.1). Step 1 transfers USDC into the wrapper's reserve account, increasing R_{USDC} . Step 2 calls `bridge_mint`, which validates that the bridge program is the authorized caller, then mints x0-USD, increasing $S_{\text{x0-USD}}$ by the same amount. The wrapper's own `bridge_mint` handler independently verifies the reserve balance, ensuring $R_{\text{USDC}}(t) \geq S_{\text{x0-USD}}(t)$ holds after every bridge mint.

15.6 Outbound Protocol (Solana → Base)

15.6.1 Step 1: Burn on Solana

The `initiate_bridge_out` instruction validates the outbound request and burns x0-USD:

The `bridge_burn` CPI performs the inverse of `bridge_mint`: it burns the sender's x0-USD tokens and transfers the equivalent USDC from the wrapper's reserve back to the bridge's reserve, preserving the wrapper reserve invariant.

Algorithm 13 Mint Execution (`execute_mint`)

```
1: procedure EXECUTEMINT(msg, proof_ctx)
2:   require msg.status = ProofVerified
3:   require proof_ctx.is_fresh()                                 $\triangleright$  Within 600s window
4:   require bridge_reserve.amount  $\geq$  msg.amount
5:
6:   // Step 1: Move USDC from bridge reserve to wrapper reserve
7:   TransferChecked(bridge_reserve  $\rightarrow$  wrapper_reserve, msg.amount)
8:
9:   // Step 2: CPI to x0-wrapper::bridge_mint
10:  x0_wrapper::bridge_mint(recipient, msg.amount)
11:
12:  msg.status  $\leftarrow$  Minted
13:  config.total_bridged_in  $+=$  msg.amount
14:  if config.total_bridged_in  $>$   $\Theta_{in}$  then
15:    config.is_paused  $\leftarrow$  true                                 $\triangleright$  Circuit breaker
16:  end if
17: end procedure
```

Algorithm 14 Outbound Bridge Initiation (`initiate_bridge_out`)

```
1: procedure BRIDGEOUT(sender, evm_recipient[20], a)
2:   require  $\neg$  config.is_paused
3:   require evm_recipient  $\neq 0^{20}$ 
4:   require a  $\in$  [MIN_AMOUNT, MAX_AMOUNT]
5:   require daily_outflow + a \leq MAX_DAILY_OUTFLOW
6:
7:   // CPI: burn x0-USD (returns USDC to bridge reserve)
8:   x0_wrapper::bridge_burn(sender, a)
9:
10:   $n \leftarrow config.bridge_out_nonce$ 
11:  config.bridge_out_nonce  $+=$  1
12:  create BridgeOutMessage(n, sender, evm_recipient, a)
13:  config.total_bridged_out  $+=$  a
14:  if config.total_bridged_out  $>$   $\Theta_{out}$  then
15:    config.is_paused  $\leftarrow$  true                                 $\triangleright$  Circuit breaker
16:  end if
17: end procedure
```

15.6.2 Step 2: Proof Generation and Unlock on Base

An off-chain relayer observes the `BridgeOutMessage` PDA, generates an SP1 STARK proof of its existence and validity, and submits the proof to the `X0UnlockContract` on Base:

Algorithm 15 Outbound Unlock (`X0UnlockContract.unlock`)

```

1: procedure UNLOCK(proofBytes, publicValues)
2:    $(P_{\text{bridge}}, n, \text{sender}, r, a, t, h) \leftarrow \text{abi.decode}(\text{publicValues})$ 
3:   require  $P_{\text{bridge}} = \text{BRIDGE\_PROGRAM\_ID}$                                  $\triangleright$  Correct program
4:   require  $\neg \text{processedNonces}[n]$                                           $\triangleright$  Replay protection
5:   require  $a \leq \text{maxPerTransaction}$ 
6:   require  $\text{dailyVolume} + a \leq \text{maxDailyVolume}$ 
7:   SP1_VERIFIER.verifyProof(programVKey, publicValues, proofBytes)
8:    $\text{processedNonces}[n] \leftarrow \text{true}$ 
9:   USDC.safeTransfer(r, a)
10:  emit BridgeUnlocked(n, sender, r, a)
11: end procedure

```

15.7 SP1 STARK Proof Circuits

Both proof circuits are implemented as SP1 guest programs that execute inside the SP1 zkVM. The zkVM produces a STARK proof that the guest program executed correctly on the claimed inputs, which can be verified on-chain in $\sim 500,000$ compute units.

15.7.1 EVM Proof Circuit (Base \rightarrow Solana)

The EVM proof circuit proves that a specific transaction was included in a finalized Base block and that its receipt contains the expected event logs.

Definition 15.6 (EVM Transaction Inclusion Proof). Let H_{rlp} denote the RLP-encoded block header, T_{rlp} the transaction, R_{rlp} the receipt, and π_T, π_R the respective MPT proofs. Given witness $w = (H_{\text{rlp}}, T_{\text{rlp}}, R_{\text{rlp}}, \pi_T, \pi_R)$, the circuit proves:

$$h_{\text{block}} = \text{keccak256}(H_{\text{rlp}}) \quad (44)$$

$$r_{\text{tx}} = \text{RLP.decode}(H_{\text{rlp}}).\text{txRoot} \quad (45)$$

$$r_{\text{rcpt}} = \text{RLP.decode}(H_{\text{rlp}}).\text{receiptsRoot} \quad (46)$$

$$\text{MPT.verify}(r_{\text{tx}}, k_T, T_{\text{rlp}}, \pi_T) = \text{true} \quad (47)$$

$$\text{MPT.verify}(r_{\text{rcpt}}, k_R, R_{\text{rlp}}, \pi_R) = \text{true} \quad (48)$$

$$R.\text{status} = 1 \quad (49)$$

and outputs public inputs $x = (h_{\text{block}}, n_{\text{block}}, h_{\text{tx}}, \text{from}, \text{to}, v, \text{success}, \text{logs})$.

The Merkle Patricia Trie (MPT) proofs anchor the transaction and receipt to the block header's state roots. Since the block hash is a public input committed to by the proof, the verifier on Solana can cross-reference it with a known finalized block hash.

```

1 pub struct EVMProofPublicInputs {
2   pub block_hash: [u8; 32],
3   pub block_number: u64,

```

```

4   pub tx_hash: [u8; 32],
5   pub from: [u8; 20],
6   pub to: [u8; 20],
7   pub value: u64,
8   pub success: bool,
9   pub event_logs: Vec<EventLog>,
10 }
11
12 pub struct EventLog {
13   pub contract_address: [u8; 20],
14   pub topics: Vec<[u8; 32]>,
15   pub data: Vec<u8>,
16 }

```

Listing 20: EVM Proof Public Inputs (sp1-evm-prover)

15.7.2 Solana Proof Circuit ($\text{Solana} \rightarrow \text{Base}$)

The Solana proof circuit proves that a `BridgeOutMessage` PDA exists in a confirmed Solana bank state, with the correct owner and data contents.

Definition 15.7 (Solana Account Inclusion Proof). Let D denote the account data, π_M the Merkle proof, B the bank hash components, and V the set of validator votes. Given witness $w = (D, \pi_M, B, V)$, the circuit proves:

$$(n, s_{\text{addr}}, r, a, t, s) = \text{parse}(D) \quad (50)$$

$$s = \text{Burned} \quad (51)$$

$$\text{owner}(D) = P_{\text{bridge}} \quad (52)$$

$$h_{\text{acct}} = \text{SHA256}(D) \quad (53)$$

$$\text{Merkle}_{16}(h_{\text{acct}}, \pi_M, \delta_{\text{hash}}) = \text{true} \quad (54)$$

$$h_{\text{bank}} = \text{SHA256}(h_{\text{parent}} \parallel \delta_{\text{hash}} \parallel \sigma_c \parallel h_{\text{blk}}) \quad (55)$$

$$\sum_{\substack{v \in V: \\ \text{Ed25519}(v.pk, h_{\text{bank}}, v.\sigma)}} v.stake \geq \frac{2}{3} \sum_{v \in V} v.stake \quad (56)$$

and outputs $x = \text{abi.encode}(P_{\text{bridge}}, n, s_{\text{addr}}, r, a, t, h_{\text{acct}})$.

Remark 15.4 (Fanout-16 Merkle Trees). Solana's accounts delta hash uses a fanout-16 Merkle tree rather than a binary tree. Each internal node hashes up to 16 children, reducing tree depth to $\log_{16} N$. The circuit verifies inclusion using this wider branching factor, with `MERKLE_FANOUT = 16` as a protocol constant.

Remark 15.5 (Validator Quorum Verification). The circuit verifies that validators representing at least $\frac{2}{3}$ of the epoch's total stake have signed the bank hash via Ed25519 signatures. This mirrors Solana's Tower BFT consensus, ensuring the proven state is finalized.

```

1 pub struct SolanaProofPublicInputs {
2   pub bridge_program_id: [u8; 32],
3   pub nonce: u64,
4   pub solana_sender: [u8; 32],
5   pub evm_recipient: [u8; 20],

```

```

6   pub amount: u64,
7   pub burn_timestamp: i64,
8   pub account_hash: [u8; 32],
9 }
10
11 impl SolanaProofPublicInputs {
12     // ABI-encodes to 7x32-byte words for EVM verification
13     pub fn abi_encode(&self) -> Vec<u8> { /* 224 bytes */ }
14 }
```

Listing 21: Solana Proof Public Inputs (sp1-solana-prover)

15.8 Rate Limiting and Circuit Breakers

FROSTGATE implements defense-in-depth rate limiting on both chains.

15.8.1 Per-Transaction Bounds

All transfer amounts are bounded:

$$\text{MIN_BRIDGE_AMOUNT} = 10 \times 10^6 \leq a \leq 100,000 \times 10^6 = \text{MAX_BRIDGE_AMOUNT_PER_TX} \quad (57)$$

corresponding to \$10–\$100,000 USDC per transaction.

15.8.2 Daily Rolling Window

Both inflow and outflow are subject to 24-hour rolling limits:

$$\sum_{\{i : t_i > t - 86400\}} a_i \leq \Phi_{\max} = 5,000,000 \times 10^6 \quad (\$5M \text{ daily}) \quad (58)$$

The daily counters reset automatically when $t_{\text{current}} - t_{\text{last_reset}} \geq 86,400$ seconds.

15.8.3 Circuit Breakers

Cumulative volume triggers an automatic pause:

$$\text{total_bridged_in} > \Theta_{\text{in}} = 10^{14} \implies \text{is_paused} \leftarrow \text{true} \quad (59)$$

The threshold $\Theta_{\text{in}} = \Theta_{\text{out}} = 100,000,000 \times 10^6$ (\$100M) provides a hard upper bound on cumulative bridge exposure. Once triggered, only the admin can unpause after investigation.

Theorem 15.2 (Bounded Bridge Exposure). *For all execution traces, the maximum USDC at risk in the bridge at any time is bounded by:*

$$\text{exposure}(t) \leq \min \left(\Theta_{\text{in}}, \Phi_{\max} \cdot \left\lceil \frac{t - t_0}{86400} \right\rceil \right) \quad (60)$$

where t_0 is the bridge deployment timestamp.

Proof. The daily inflow cap ensures no more than Φ_{\max} enters per 24-hour window. The circuit breaker provides an absolute ceiling of Θ_{in} . Since locked USDC on Base can only be unlocked

via a valid outbound proof, the exposure is bounded by the lesser of the cumulative cap and the sum of daily windows.

□

□

15.9 Proof Freshness

To prevent indefinite replay of verified proofs, proof contexts enforce a temporal validity window:

Definition 15.8 (Bridge Proof Freshness). An `EVMProofContext` C is *fresh* at time t if and only if:

$$t - C.\text{verified_at} < \Delta_{\text{bridge}} = 600 \text{ seconds} \quad (61)$$

The 10-minute window accounts for Hyperlane relay latency ($\sim 2\text{-}3$ minutes), SP1 proving time (~ 1 minute), and transaction submission buffer. The window is deliberately larger than the ZK verifier's 300-second window (Section 9) to accommodate cross-chain timing uncertainty.

15.10 EVM Contracts

Both Solidity contracts inherit OpenZeppelin's `Ownable`, `Pausable`, and `ReentrancyGuard`, providing:

- **Pausability:** Owner can halt all lock/unlock operations in an emergency.
- **Reentrancy protection:** Guards against callback-based exploits during USDC transfers.
- **Rate limiting:** Independent daily volume caps and per-transaction limits on both contracts.
- **Admin recovery:** Timelocked `adminUnlock` on the lock contract allows fund recovery if the Solana-side flow permanently fails.

15.10.1 XOLockContract

The lock contract manages inbound flow (Base → Solana):

1. Receives USDC via `SafeERC20.safeTransferFrom`
2. Validates amount bounds and daily rate limits
3. Dispatches the 80-byte message body via Hyperlane's `IMailbox.dispatch`
4. Emits the `Locked` event with the sender, recipient, amount, nonce, and Hyperlane message ID
5. Increments the per-user nonce for replay ordering

15.10.2 XOUnlockContract

The unlock contract manages outbound flow (Solana → Base):

1. Receives an SP1 proof and ABI-encoded public values
2. Verifies the proof via `ISP1Verifier.verifyProof(vkey, values, proof)`

3. Validates the bridge program ID matches the expected Solana program
4. Checks replay protection via the `processedNonces` mapping
5. Transfers USDC to the EVM recipient via `SafeERC20.safeTransfer`

15.11 Administrative Controls

FROSTGATE administrative operations are subject to a 48-hour timelock, matching the x0-wrapper governance model (Section 8):

1. `schedule_admin_action`: Creates a `BridgeAdminAction` PDA with a monotonic nonce and records the action type (e.g., add/remove EVM contract, add/remove domain, change trusted program addresses).
2. `execute_admin_action`: Executes after `current_time-scheduled_at` $\geq 172,800$ seconds. Validates the action nonce, type, and that the action has not been cancelled.
3. `cancel_admin_action`: Allows the admin to abort a scheduled action before execution.

This timelock gives users and monitoring systems 48 hours to detect and respond to potentially malicious configuration changes (e.g., adding an attacker-controlled EVM contract to the whitelist).

15.12 Security Analysis

15.12.1 Threat Model

We extend the base protocol's threat model (Section 10) with bridge-specific adversaries:

1. **Proof Forgery**: Adversary attempts to fabricate an SP1 proof for a nonexistent deposit.
2. **Amount Inflation**: Adversary deposits a but attempts to mint $a' > a$ on the destination chain.
3. **Recipient Redirection**: Adversary deposits for recipient r but attempts to mint for $r' \neq r$.
4. **Replay Attack**: Adversary attempts to reuse a valid proof to mint twice.
5. **Hyperlane Compromise**: Adversary controls the Hyperlane relayer and crafts arbitrary messages.
6. **EVM Contract Impersonation**: Adversary deploys a malicious contract that emits fake `Locked` events.

15.12.2 Mitigations

Theorem 15.3 (FROSTGATE Soundness). *Under the soundness of the SP1 STARK proof system, the collision resistance of SHA-256 and Keccak-256, and the unforgeability of Ed25519 signatures, FROSTGATE is a sound cross-chain bridge (Definition 14.2).*

Attack Vector	Mitigation
Proof Forgery	SP1 STARK soundness ($\leq 2^{-100}$)
Amount Inflation	Event amount validated against BridgeMessage
Recipient Redirect	Event recipient validated against BridgeMessage
Replay (inbound)	BridgeMessage PDA uniqueness + status machine
Replay (outbound)	<code>processedNonces</code> mapping on EVM
Hyperlane Compromise	SP1 proof required <i>in addition</i> to Hyperlane message
Contract Impersonation	Event contract address checked against whitelist

Table 5: FROSTGATE Attack Vectors and Mitigations

Proof. We show that no minting can occur without a corresponding genuine deposit.

For *inbound* transfers (Base → Solana):

1. `execute_mint` requires `msg.status = ProofVerified`.
2. The `ProofVerified` status is only set by `verify_evm_proof`, which validates an SP1 STARK proof.
3. The STARK proof commits to a block hash, transaction hash, and event logs (Definition 15.6).
4. `validate_locked_event` (Theorem 15.1) ensures the proven event logs contain a genuine `Locked` event with matching amount, recipient, and nonce from a whitelisted contract.
5. Therefore, minting requires a valid EVM transaction that deposited the exact amount for the exact recipient.

For *outbound* transfers (Solana → Base):

1. `XOUnlockContract.unlock` requires a valid SP1 proof of a `BridgeOutMessage` PDA.
2. The STARK proof commits to the account hash, nonce, amount, and recipient (Definition 15.7).
3. Validator quorum verification (Equation 56) ensures the proven state is finalized under Solana’s Tower BFT consensus.
4. The `processedNonces` mapping prevents replay.
5. Therefore, unlocking requires a genuine burn transaction on Solana with finalized consensus. \square

\square

Theorem 15.4 (FROSTGATE Completeness). *FROSTGATE is a complete cross-chain bridge (Definition 14.3), given liveness of the Hyperlane relayer (inbound) and the SP1 prover operator (both directions).*

Proof. For any valid `Lock` on Base: (1) the Hyperlane relayer delivers the message, creating a `BridgeMessage` PDA; (2) the SP1 prover generates a valid proof of the EVM transaction; (3) the `verify_evm_proof` instruction verifies the proof and transitions the message to `ProofVerified`;

(4) the `execute_mint` instruction mints x0-USD. Each step is deterministic given valid inputs. The only liveness assumptions are on the relayer and prover, which are permissionless roles. \square

\square

Theorem 15.5 (Reserve Invariant Under Bridge Operations). *For all execution traces involving bridge mint and burn operations:*

$$\forall t : R_{USDC}(t) \geq S_{x0\text{-USD}}(t) \quad (62)$$

Proof. Bridge mint: `execute_mint` first transfers a USDC from the bridge reserve to the wrapper reserve ($R_{USDC} += a$), then calls `bridge_mint` to mint a x0-USD ($S_{x0\text{-USD}} += a$). The wrapper's `bridge_mint` handler independently verifies $R_{USDC} \geq S_{x0\text{-USD}}$ before minting.

Bridge burn: `bridge_burn` burns a x0-USD ($S_{x0\text{-USD}} -= a$) and transfers a USDC from the wrapper reserve to the bridge reserve ($R_{USDC} -= a$). The decrement is symmetric.

In both cases the invariant is preserved by the wrapper's own validation (Theorem 8.1). \square

\square

16 Future Work

16.1 Extended Zero-Knowledge Applications

With the foundational ZK verification infrastructure now deployed (Section 9), several extensions become feasible:

- **Private spend limit verification:** Prove that a transfer is within policy limits without revealing the daily limit or cumulative spend, using range proofs composed with the existing Groth16 circuits
- **Proof of reputation threshold:** Prove $S \geq \theta$ for a reputation score S and threshold θ without revealing the exact score, enabling privacy-preserving agent discovery
- **Confidential escrow amounts:** Extend the escrow state machine to operate on encrypted amounts, using ZK proofs to verify state transitions without revealing the escrowed value
- **Recursive proof composition:** Batch multiple proof verifications into a single on-chain instruction using recursive SNARKs, reducing per-transaction compute costs

16.2 Cross-Chain Extension

With FROSTGATE operational on Base \longleftrightarrow Solana (Section 14), the architecture generalizes to additional EVM chains:

- **Multi-chain deployment:** Deploy `XOLockContract` and `XOUnlockContract` instances on Ethereum mainnet, Arbitrum, Optimism, and other EVM chains, each with its own Hyperlane domain ID registered in `BridgeConfig`
- **Non-EVM chains:** Extend the SP1 proof model to verify state inclusion on non-EVM chains (e.g., Cosmos via IBC, Aptos via Move proofs), enabling a unified cross-chain payment mesh

- **Unified reputation across chains:** Anchor reputation scores to Solana as the settlement layer, with bridge-proven attestations propagating trust scores cross-chain
- **Cross-chain escrow settlement:** Compose FROSTGATE with the escrow protocol (Section 5) to enable trustless cross-chain escrow with atomic settlement guarantees

16.3 Machine Learning Integration

- Anomaly detection for suspicious agent behavior using on-chain transaction patterns
- Dynamic reputation weighting via learned feature importance
- Predictive escrow dispute resolution using historical outcome data

16.4 Governance

- DAO for protocol parameter updates (fee rates, timelock durations)
- Community-driven arbiter selection and staking
- Fee redistribution to token holders and active arbiters

17 Conclusion

We have presented x0, a comprehensive payment infrastructure for autonomous agents. The protocol’s key innovations include:

1. **Transfer Hook Policy Enforcement:** Programmable spending limits enforced cryptographically on every transaction
2. **HTTP 402 Protocol:** Standardized payment negotiation for agent-to-agent commerce
3. **Conditional Escrow:** Trustless high-value transactions with dispute resolution
4. **Temporal Reputation:** On-chain trust scores with decay to prevent stale reputations
5. **USDC Wrapper:** Stable payments with cryptographic reserve invariants
6. **Human-in-the-Loop:** Fallback mechanism for exceptional cases
7. **Zero-Knowledge Verification:** On-chain Groth16 proof verification enabling confidential transfers with provable correctness
8. **FROSTGATE Bridge:** Trustless bidirectional Base \longleftrightarrow Solana bridging using SP1 STARK proofs and Hyperlane message passing, with event-level validation, rate limiting, and circuit breakers—eliminating all off-chain trust assumptions from cross-chain agent payments

The protocol has been deployed on Solana devnet and is undergoing security audits. We invite the community to review the codebase at github.com/x0-protocol and provide feedback.

The future of AI agent commerce requires infrastructure that is:

- **Fast:** Sub-second transaction finality
- **Cheap:** Sub-cent transaction costs
- **Programmable:** Enforcing complex spending rules
- **Trustless:** No centralized intermediaries
- **Safe:** Bounded downside risk

x0 achieves these properties through careful cryptographic design and integration with Solana's high-performance blockchain.

Acknowledgments

References

- [1] Nakamoto, S. (2008). Bitcoin: A Peer-to-Peer Electronic Cash System.
- [2] Buterin, V. (2014). Ethereum: A Next-Generation Smart Contract and Decentralized Application Platform.
- [3] Yakovenko, A. (2018). Solana: A new architecture for a high performance blockchain.
- [4] Poon, J., & Dryja, T. (2016). The Bitcoin Lightning Network: Scalable Off-Chain Instant Payments.
- [5] Raiden Network Team. (2017). Raiden Network: Fast, cheap, scalable token transfers for Ethereum.
- [6] Solana Labs. (2023). Token-2022 Extension Documentation. <https://spl.solana.com/token-2022>
- [7] ElGamal, T. (1985). A public key cryptosystem and a signature scheme based on discrete logarithms. IEEE Transactions on Information Theory, 31(4), 469-472.
- [8] Merkle, R. C. (1988). A Digital Signature Based on a Conventional Encryption Function. CRYPTO '87.
- [9] Bloom, B. H. (1970). Space/time trade-offs in hash coding with allowable errors. Communications of the ACM, 13(7), 422-426.
- [10] Bernstein, D. J., et al. (2012). High-speed high-security signatures. Journal of Cryptographic Engineering, 2(2), 77-89.
- [11] Groth, J. (2016). On the Size of Pairing-Based Non-interactive Arguments. EUROCRYPT 2016, Lecture Notes in Computer Science, vol 9666, pp. 305-326.
- [12] Hamburg, M. (2015). Decaf: Eliminating cofactors through point compression. CRYPTO 2015. Extended to Ristretto255 for use with Curve25519.

- [13] Succinct Labs. (2024). SP1: A performant, open-source zero-knowledge virtual machine. <https://docs.succinct.xyz>
- [14] Hyperlane. (2024). Hyperlane: Permissionless Interoperability. <https://docs.hyperlane.xyz>
- [15] Ben-Sasson, E., Bentov, I., Horesh, Y., & Riabzev, M. (2018). Scalable, transparent, and post-quantum secure computational integrity. IACR Cryptology ePrint Archive, 2018, 046.
- [16] Wood, G. (2014). Ethereum: A Secure Decentralised Generalised Transaction Ledger. Ethereum Yellow Paper, Appendix D: Modified Merkle Patricia Trie.

A Error Codes Reference

Error codes follow a structured numbering scheme where the high byte identifies the program category and the low nibble groups errors by subcategory:

Range	Name	Description
<i>x0-guard: Policy Errors (0x1100–0x110F)</i>		
0x1101	RecipientNotWhitelisted	Recipient not in whitelist
0x1102	DailyLimitExceeded	Rolling 24h limit exceeded
0x1103	InvalidMerkleProof	Merkle proof verification failed
0x1104	InvalidBloomFilter	Bloom filter misconfigured
0x1105	PolicyNotFound	Agent policy PDA missing
0x1106	UnauthorizedSigner	Unauthorized transaction signer
0x1107	ConfidentialTransferFailed	ZK proof verification failed
<i>x0-guard: Configuration Errors (0x1110–0x111F)</i>		
0x1110	DailyLimitTooHigh	Exceeds maximum daily limit
0x1111	DailyLimitTooLow	Below minimum daily limit
0x1112	InvalidWhitelistConfig	Invalid whitelist configuration
0x1118	DelegationRequired	Agent must be delegate, not owner
0x111B	BoundTokenAccountMismatch	Wrong source token account
0x111F	MissingZkProof	ZK proof required for confidential transfer
<i>x0-guard: Transfer Errors (0x1120–0x112F)</i>		
0x1120	ZeroTransferAmount	Transfer amount is zero
0x1121	TransferAmountTooHigh	Exceeds per-transaction limit
<i>x0-guard: Blink Errors (0x1130–0x113F)</i>		
0x1130	BlinkRateLimitExceeded	>3 blinks per hour
0x1131	BlinkExpired	Blink has expired
<i>x0-guard: Severity Errors (0x1140–0x114F)</i>		
0x1140	PolicyUpdateTooFrequent	Rate-limited policy update
0x1141	SingleTransactionLimitExceeded	Per-tx limit exceeded
0x1143	ExtraMetasAlreadyInitialized	Extra metas re-init blocked
0x1144	UnauthorizedExtraMetasInitializer	Mint authority required
<i>x0-escrow: State Errors (0x1108–0x1109, 0x1200+)</i>		
0x1108	EscrowExpired	Escrow timeout reached
0x1109	InvalidEscrowState	Invalid state transition
0x1202	SameBuyerAndSeller	Buyer and seller must differ
0x1204	ZeroEscrowAmount	Escrow amount is zero
<i>x0-escrow: Operation Errors (0x1210–0x121F)</i>		
0x1210	OnlyBuyerCanFund	Only buyer can fund
0x1211	OnlySellerCanDeliver	Only seller can mark delivered
0x1213	OnlyArbiterCanResolve	Only arbiter can resolve
0x1217	ArbiterResolutionTooEarly	Arbiter must wait for delay
<i>x0-registry: Entry Errors (0x1300–0x130F)</i>		
0x1300	AgentAlreadyRegistered	Agent already registered
0x1301	AgentNotFound	Agent not found
0x1303	EndpointTooLong	Endpoint URL too long
0x1304	TooManyCapabilities	Exceeds max capabilities
0x1307	InsufficientListingFee	Listing fee not met
<i>x0-reputation: Errors (0x1400–0x140F)</i>		

Range	Name	Description
0x1400	ReputationNotFound	Reputation account missing
0x1401	ReputationAlreadyInitialized	Already initialized
0x1403	UnauthorizedReputationUpdate	Unauthorized caller
0x1404	InsufficientTransactions	Too few txns for score
<i>x0-token: Token Errors (0x1500–0x1515)</i>		
0x1500	MintAlreadyInitialized	Mint already initialized
0x1503	ConfidentialTransfersNotEnabled	Confidential not enabled
0x1507	InvalidElGamalPubkey	Invalid ElGamal public key
0x1508	InvalidPubkeyValidityProof	Invalid pubkey proof
0x150A	InvalidWithdrawProof	Invalid withdraw proof
0x150F	AmountExceedsConfidentialMax	Amount > $2^{48} - 1$
<i>x0-wrapper: State/Reserve Errors (0x1600–0x164F)</i>		
0x1600	WrapperPaused	Operations are paused
0x1610	InsufficientReserve	Insufficient reserve balance
0x1611	ReserveInvariantViolated	Reserve < supply
0x1620	DepositTooSmall	Below minimum deposit
0x1623	DailyRedemptionLimitExceeded	Daily limit exceeded
0x1630	Unauthorized	Admin required
0x1640	AdminActionNotFound	Action PDA not found
0x1643	TimelockNotExpired	Timelock period pending
<i>x0-zk-verifier: Proof Errors (0x1700–0x1721)</i>		
0x1700	ProofVerificationFailed	ZK proof verification failed
0x1701	InvalidProofData	Invalid proof data format
0x1703	InvalidProofType	Invalid proof type
0x1704	ProofExpired	Proof timestamp too old
0x1710	AmountTooLarge	Amount > $2^{48} - 1$
0x1711	InvalidElGamalPubkey	Invalid ElGamal public key
0x1713	ProofSizeMismatch	Proof data size mismatch
0x1720	ArithmeticOverflow	Arithmetic overflow
<i>x0-bridge: Configuration Errors (0x1800–0x1806)</i>		
0x1800	BridgeAlreadyInitialized	Bridge already initialized
0x1802	BridgePaused	Bridge operations are paused
0x1803	Unauthorized	Admin required
0x1805	InvalidSP1Verifier	Invalid SP1 verifier program
<i>x0-bridge: Message Errors (0x1810–0x1819)</i>		
0x1810	UnsupportedDomain	Origin domain not supported
0x1811	UnauthorizedSenderContract	Sender contract not in allowed list
0x1812	MessageAlreadyProcessed	Message already processed
0x1817	InvalidProcessAuthority	PDA derivation mismatch
0x1818	InvalidSenderFormat	EVM sender zero-padding check
0x1819	CircuitBreakerTriggered	Bridge volume exceeds threshold
<i>x0-bridge: Proof Errors (0x1820–0x182D)</i>		
0x1820	ProofVerificationFailed	STARK proof verification failed
0x1823	ProofExpired	Proof exceeded validity window
0x1826	DepositEventNotFound	Locked event not in proof logs
0x1828	EventContractMismatch	Event from unauthorized contract
0x1829	EventRecipientMismatch	Event recipient ≠ message recipient
0x182A	EventAmountMismatch	Event amount ≠ message amount
0x182B	EventNonceMismatch	Event nonce ≠ expected nonce
<i>x0-bridge: Rate Limiting (0x1830–0x183F)</i>		
0x1830	AmountTooSmall	Bridge amount < minimum
0x1831	AmountTooLarge	Bridge amount > per-tx max
0x1832	DailyInflowLimitExceeded	Daily inflow limit exceeded
0x1833	InsufficientBridgeReserve	Insufficient reserve for minting
<i>x0-bridge: Outbound Errors (0x1870–0x1876)</i>		
0x1870	BridgeOutPaused	Outbound operations paused
0x1873	DailyOutflowLimitExceeded	Daily outflow limit exceeded
0x1874	OutboundCircuitBreakerTriggered	Outbound circuit breaker
0x1875	InvalidEvmRecipient	Invalid EVM recipient address

Table 6: Protocol Error Codes (representative subset; 8 enums, 90+ total codes). Full reference in source: [x0-common/src/error.rs](#)

Program	Devnet Address
x0-guard	2uYGW3fQUGfhrwVbkupdasXBpRPfGYBGTLUdaPTXU9vP
x0-token	EHHTCSyGkmnsBhGsvCmLzKgcSxtsN31ScrfiwcCbjHci
x0-escrow	AhaDyVm8LBxpUwFdArA37LnHvNx6cNWe3KAiy8zGqhHF
x0-registry	Bebty49EPhFoANKDw7TqLQ2bX61ackNav5iNkj36eVJo
x0-reputation	FfzkTWRGAJQPDePbjZdEhKHqC1UpqvDrpv4TEiWpx6y
x0-wrapper	EomiXBbg94Smu4ipDoJtuguazcd1KjLFDFJt2fCabvJ8
x0-zk-verifier	zQWSrznKgcK8aHA4ry7xbSCdP36FqgUHj766YM3pwre
x0-bridge	4FuyKfQysHxcTeNJtz5rBzzS8kmjn2DdkgXH1Q7edXa7

Table 7: Deployed Program Addresses

B Program Addresses

C Mathematical Notation

Symbol	Meaning
H	Cryptographic hash function (SHA-256)
sk, pk	Secret key, public key
σ	Digital signature
L	Daily spending limit
x_i	Transaction amount at index i
t_i	Transaction timestamp at index i
R	Recipient address
W	Whitelist set
n	Number of transactions
m	Number of whitelist entries
k	Number of hash functions (Bloom)
S	Reputation score
ρ	Reserve ratio
R_{USDC}	USDC reserve balance
$S_{\text{x0-USD}}$	x0-USD supply
G, H	Ristretto255 generators
P	ElGamal public key ($P = s \cdot G$)
s	ElGamal secret scalar
π	Zero-knowledge proof ($\pi = (A, B, C)$ for Groth16)
Δ_{proof}	Proof freshness window (300 seconds)
Δ_{bridge}	Bridge proof freshness window (600 seconds)
$\Theta_{\text{in}}, \Theta_{\text{out}}$	Circuit breaker thresholds (\$100M)
Φ_{max}	Maximum daily bridge volume (\$5M)
δ_{hash}	Solana accounts delta hash
h_{bank}	Solana bank hash

Table 8: Mathematical Notation Reference

D SDK Integration Guide

D.1 Installation

The x0 TypeScript SDK is available via npm:

```

1 npm install @x0-protocol/sdk
2 # or
3 yarn add @x0-protocol/sdk

```

Listing 22: SDK Installation

D.2 Quick Start

```

1 import {
2   X0Client,
3   EscrowClient,
4   ReputationClient,
5   createPaymentRequest,
6   verifyPaymentProof
7 } from '@x0-protocol/sdk';
8 import { Connection, Keypair } from '@solana/web3.js';
9
10 // Initialize connection
11 const connection = new Connection('https://api.devnet.solana.com');
12 const wallet = Keypair.generate(); // or load from file
13
14 // Create escrow
15 const escrow = new EscrowClient(connection);
16 const { escrowPda, tx } = await escrow.buildCreateEscrowInstruction({
17   buyer: wallet.publicKey,
18   seller: sellerPubkey,
19   amount: 1_000_000, // 1 USDC (6 decimals)
20   memo: 'API service payment',
21   timeoutSeconds: 86400, // 24 hours
22 });
23
24 // x402 payment flow (service side)
25 const paymentRequest = createPaymentRequest({
26   recipient: servicePubkey,
27   amount: '1000000',
28   resource: '/api/v1/generate',
29   network: 'solana-devnet',
30 });
31
32 // Verify payment (service side)
33 const isValid = await verifyPaymentProof(connection, proof, {
34   expectedRecipient: servicePubkey,
35   expectedAmount: 1_000_000,
36   toleranceSeconds: 300,
37 });

```

Listing 23: Basic SDK Usage

D.3 SDK Modules

D.4 Resources

- **npm:** <https://npmjs.com/package/@x0-protocol/sdk>
- **GitHub:** <https://github.com/0xtnxl/x0/sdk>

Module	Purpose
EscrowClient	Create, fund, release, dispute escrows
ReputationClient	Query and update agent reputation
RegistryClient	Register agents, discover services
GuardClient	Manage spending policies
WrapperClient	Wrap/unwrap USDC to x0-USD
ConfidentialClient	Confidential transfer proof generation via WASM
ZkVerifierClient	On-chain ZK proof verification
x402	HTTP 402 payment request/proof utilities
blink	Generate Solana Actions for human approval
BridgeClient	Cross-chain bridge operations (lock, verify, mint, bridge-out)

Table 9: SDK Module Overview

- **API Docs:** <https://docs.x0protocol.dev/sdk>
- **Examples:** <https://github.com/0xtnxl/x0/examples>

D.5 Confidential Transfer Architecture

For the complete treatment of the zero-knowledge proof system—including cryptographic foundations, proof types, the `ProofContext` state account, and the end-to-end verification flow—see Section 9. The SDK’s `ConfidentialClient` module wraps the WASM-compiled `x0-zk-proofs` crate to provide transparent proof generation:

```

1 import { ConfidentialClient } from '@x0-protocol/sdk';
2
3 const confidentialClient = new ConfidentialClient(connection, wallet);
4
5 // Proof generation via WASM (x0-zk-proofs):
6 // - PubkeyValidityProof: Proves ElGamal key validity (64 bytes)
7 // - WithdrawProof: Proves withdrawal against encrypted balance (160 bytes)
8 // - ZeroBalanceProof: Proves zero balance for account closure (96 bytes)
9 // On-chain verification via x0-zk-verifier creates ProofContext PDAs
10 // that Token-2022 references during confidential transfer execution.
```

Listing 24: Confidential Transfer ZK Proof Generation